



UPPSALA
UNIVERSITET

Bayesian Statistics and Data Analysis

Lecture 5

Måns Magnusson

Department of Statistics, Uppsala University
Thanks to Aki Vehtari, Aalto University

- Monte Carlo recap
- Markov Chain Monte Carlo (MCMC)
 - Gibbs sampling
 - Metropolis-Hastings
- MCMC Diagnostics
 - Warm-up
 - Convergence
 - S_{eff} , MCSE, and autocorrelation
- Difficult geometries
- Additional material



UPPSALA
UNIVERSITET

Course announcements

- Monte Carlo recap
 - Markov Chain Monte Carlo (MCMC)
 - Gibbs sampling
 - Metropolis-Hastings
 - MCMC Diagnostics
 - Warm-up
 - Convergence
 - S_{eff} , MCSE, and autocorrelation
 - Difficult geometries
 - Additional material
- Oral exam: what worked well and what should change?
 - Mini-project proposals



UPPSALA
UNIVERSITET

It's all about expectations

$$E_{p(\theta|y)}[f(\theta)] = \int f(\theta) p(\theta|y) d\theta,$$

$$\text{where } p(\theta|y) = \frac{p(y|\theta)p(\theta)}{\int p(y|\theta)p(\theta)d\theta}$$

- Monte Carlo recap
- Markov Chain Monte Carlo (MCMC)
 - Gibbs sampling
 - Metropolis-Hastings
- MCMC Diagnostics
 - Warm-up
 - Convergence
 - S_{eff} , MCSE, and autocorrelation
- Difficult geometries
- Additional material



It's all about expectations

$$E_{p(\theta|y)}[f(\theta)] = \int f(\theta) p(\theta|y) d\theta,$$

$$\text{where } p(\theta|y) = \frac{p(y|\theta)p(\theta)}{\int p(y|\theta)p(\theta)d\theta}$$

We can easily evaluate $p(y|\theta)p(\theta)$ for any θ , but the integral $\int p(y|\theta)p(\theta)d\theta$ is usually difficult.

- Monte Carlo recap
- Markov Chain Monte Carlo (MCMC)
 - Gibbs sampling
 - Metropolis-Hastings
- MCMC Diagnostics
 - Warm-up
 - Convergence
 - S_{eff} , MCSE, and autocorrelation
- Difficult geometries
- Additional material



It's all about expectations

$$E_{p(\theta|y)}[f(\theta)] = \int f(\theta) p(\theta|y) d\theta,$$

$$\text{where } p(\theta|y) = \frac{p(y|\theta)p(\theta)}{\int p(y|\theta)p(\theta)d\theta}$$

We can easily evaluate $p(y|\theta)p(\theta)$ for any θ , but the integral $\int p(y|\theta)p(\theta)d\theta$ is usually difficult.

Define the unnormalized posterior $q(\theta) = p(y|\theta)p(\theta)$, so that $p(\theta|y) \propto q(\theta)$. We can use q , for example, in

- Monte Carlo recap
- Markov Chain Monte Carlo (MCMC)
 - Gibbs sampling
 - Metropolis-Hastings
- MCMC Diagnostics
 - Warm-up
 - Convergence
 - S_{eff} , MCSE, and autocorrelation
- Difficult geometries
- Additional material



It's all about expectations

$$E_{p(\theta|y)}[f(\theta)] = \int f(\theta)p(\theta|y)d\theta,$$

$$\text{where } p(\theta|y) = \frac{p(y|\theta)p(\theta)}{\int p(y|\theta)p(\theta)d\theta}$$

We can easily evaluate $p(y|\theta)p(\theta)$ for any θ , but the integral $\int p(y|\theta)p(\theta)d\theta$ is usually difficult.

Define the unnormalized posterior $q(\theta) = p(y|\theta)p(\theta)$, so that $p(\theta|y) \propto q(\theta)$. We can use q , for example, in

- Monte Carlo methods which can sample from $p(\theta^{(s)}|y)$ using only $q(\theta^{(s)})$

$$E_{p(\theta|y)}[f(\theta)] \approx \frac{1}{S} \sum_{s=1}^S f(\theta^{(s)})$$

- Monte Carlo recap
- Markov Chain Monte Carlo (MCMC)
 - Gibbs sampling
 - Metropolis-Hastings
- MCMC Diagnostics
 - Warm-up
 - Convergence
 - S_{eff} , MCSE, and autocorrelation
- Difficult geometries
- Additional material



UPPSALA
UNIVERSITET

Monte Carlo

- Monte Carlo recap
 - Markov Chain Monte Carlo (MCMC)
 - Gibbs sampling
 - Metropolis-Hastings
 - MCMC Diagnostics
 - Warm-up
 - Convergence
 - S_{eff} , MCSE, and autocorrelation
 - Difficult geometries
 - Additional material
- Monte Carlo methods we have discussed so far
 - Inverse CDF works mainly for 1D



UPPSALA
UNIVERSITET

Monte Carlo

- Monte Carlo recap
 - Markov Chain Monte Carlo (MCMC)
 - Gibbs sampling
 - Metropolis-Hastings
 - MCMC Diagnostics
 - Warm-up
 - Convergence
 - S_{eff} , MCSE, and autocorrelation
 - Difficult geometries
 - Additional material
- Monte Carlo methods we have discussed so far
 - Inverse CDF works mainly for 1D
 - Grid methods work in a few dimensions



- Monte Carlo recap
 - Markov Chain Monte Carlo (MCMC)
 - Gibbs sampling
 - Metropolis-Hastings
 - MCMC Diagnostics
 - Warm-up
 - Convergence
 - S_{eff} , MCSE, and autocorrelation
 - Difficult geometries
 - Additional material
- Monte Carlo methods we have discussed so far
 - Inverse CDF works mainly for 1D
 - Grid methods work in a few dimensions
 - Rejection sampling typically becomes inefficient as dimension increases (truncation is a special case)



- Monte Carlo recap
- Markov Chain Monte Carlo (MCMC)
 - Gibbs sampling
 - Metropolis-Hastings
- MCMC Diagnostics
 - Warm-up
 - Convergence
 - S_{eff} , MCSE, and autocorrelation
- Difficult geometries
- Additional material
- Monte Carlo methods we have discussed so far
 - Inverse CDF works mainly for 1D
 - Grid methods work in a few dimensions
 - Rejection sampling typically becomes inefficient as dimension increases (truncation is a special case)
 - Importance sampling is reliable only in special cases



- Monte Carlo recap
 - Markov Chain Monte Carlo (MCMC)
 - Gibbs sampling
 - Metropolis-Hastings
 - MCMC Diagnostics
 - Warm-up
 - Convergence
 - S_{eff} , MCSE, and autocorrelation
 - Difficult geometries
 - Additional material
- Monte Carlo methods we have discussed so far
 - Inverse CDF works mainly for 1D
 - Grid methods work in a few dimensions
 - Rejection sampling typically becomes inefficient as dimension increases (truncation is a special case)
 - Importance sampling is reliable only in special cases
 - What to do in high dimensions?



- Monte Carlo recap
 - Markov Chain Monte Carlo (MCMC)
 - Gibbs sampling
 - Metropolis-Hastings
 - MCMC Diagnostics
 - Warm-up
 - Convergence
 - S_{eff} , MCSE, and autocorrelation
 - Difficult geometries
 - Additional material
- Monte Carlo methods we have discussed so far
 - Inverse CDF works mainly for 1D
 - Grid methods work in a few dimensions
 - Rejection sampling typically becomes inefficient as dimension increases (truncation is a special case)
 - Importance sampling is reliable only in special cases
 - What to do in high dimensions?
 - Markov chain Monte Carlo (Ch 11-12)



- Monte Carlo recap
 - Markov Chain Monte Carlo (MCMC)
 - Gibbs sampling
 - Metropolis-Hastings
 - MCMC Diagnostics
 - Warm-up
 - Convergence
 - S_{eff} , MCSE, and autocorrelation
 - Difficult geometries
 - Additional material
- Monte Carlo methods we have discussed so far
 - Inverse CDF works mainly for 1D
 - Grid methods work in a few dimensions
 - Rejection sampling typically becomes inefficient as dimension increases (truncation is a special case)
 - Importance sampling is reliable only in special cases
 - What to do in high dimensions?
 - Markov chain Monte Carlo (Ch 11-12)
 - Laplace, Variational*, EP* (Ch 4, 13*, next course)



Markov chains

- A Markov chain is a sequence in which the next state depends on the current state, but not on earlier states

$$p(\theta^{(t)} \mid \theta^{(t-1)}, \dots, \theta^{(0)}) = p(\theta^{(t)} \mid \theta^{(t-1)}).$$

- Monte Carlo recap
- Markov Chain Monte Carlo (MCMC)
 - Gibbs sampling
 - Metropolis-Hastings
- MCMC Diagnostics
 - Warm-up
 - Convergence
 - S_{eff} , MCSE, and autocorrelation
- Difficult geometries
- Additional material



Markov chains

- A Markov chain is a sequence in which the next state depends on the current state, but not on earlier states

$$p(\theta^{(t)} \mid \theta^{(t-1)}, \dots, \theta^{(0)}) = p(\theta^{(t)} \mid \theta^{(t-1)}).$$

- $T(\theta^{(t)} \mid \theta^{(t-1)})$ is the **transition distribution**

- Monte Carlo recap
- **Markov Chain Monte Carlo (MCMC)**
 - Gibbs sampling
 - Metropolis-Hastings
- MCMC Diagnostics
 - Warm-up
 - Convergence
 - S_{eff} , MCSE, and autocorrelation
- Difficult geometries
- Additional material



Markov chains

- A Markov chain is a sequence in which the next state depends on the current state, but not on earlier states

$$p(\theta^{(t)} \mid \theta^{(t-1)}, \dots, \theta^{(0)}) = p(\theta^{(t)} \mid \theta^{(t-1)}).$$

- $T(\theta^{(t)} \mid \theta^{(t-1)})$ is the **transition distribution**
- For a well-behaved chain, the distribution P_t of $\theta^{(t)}$ approaches a stationary distribution π

- Monte Carlo recap
- **Markov Chain Monte Carlo (MCMC)**
 - Gibbs sampling
 - Metropolis-Hastings
- MCMC Diagnostics
 - Warm-up
 - Convergence
 - S_{eff} , MCSE, and autocorrelation
- Difficult geometries
- Additional material



Markov chains

- A Markov chain is a sequence in which the next state depends on the current state, but not on earlier states

$$p(\theta^{(t)} \mid \theta^{(t-1)}, \dots, \theta^{(0)}) = p(\theta^{(t)} \mid \theta^{(t-1)}).$$

- $T(\theta^{(t)} \mid \theta^{(t-1)})$ is the **transition distribution**
- For a well-behaved chain, the distribution P_t of $\theta^{(t)}$ approaches a stationary distribution π
- Total variation distance is

$$d_{\text{TV}}(P_t, \pi) = \sup_A |P_t(A) - \pi(A)|.$$

It is the largest difference between the probabilities that P_t and π assign to the same event A

- Monte Carlo recap
- Markov Chain Monte Carlo (MCMC)
 - Gibbs sampling
 - Metropolis-Hastings
- MCMC Diagnostics
 - Warm-up
 - Convergence
 - S_{eff} , MCSE, and autocorrelation
- Difficult geometries
- Additional material



- A Markov chain is a sequence in which the next state depends on the current state, but not on earlier states

$$p(\theta^{(t)} \mid \theta^{(t-1)}, \dots, \theta^{(0)}) = p(\theta^{(t)} \mid \theta^{(t-1)}).$$

- $T(\theta^{(t)} \mid \theta^{(t-1)})$ is the **transition distribution**
- For a well-behaved chain, the distribution P_t of $\theta^{(t)}$ approaches a stationary distribution π
- Total variation distance is

$$d_{\text{TV}}(P_t, \pi) = \sup_A |P_t(A) - \pi(A)|.$$

It is the largest difference between the probabilities that P_t and π assign to the same event A

- Convergence in total variation means $d_{\text{TV}}(P_t, \pi) \rightarrow 0$

- Monte Carlo recap
- Markov Chain Monte Carlo (MCMC)
 - Gibbs sampling
 - Metropolis-Hastings
- MCMC Diagnostics
 - Warm-up
 - Convergence
 - S_{eff} , MCSE, and autocorrelation
- Difficult geometries
- Additional material



Markov chains

- A Markov chain is a sequence in which the next state depends on the current state, but not on earlier states

$$p(\theta^{(t)} \mid \theta^{(t-1)}, \dots, \theta^{(0)}) = p(\theta^{(t)} \mid \theta^{(t-1)}).$$

- $T(\theta^{(t)} \mid \theta^{(t-1)})$ is the **transition distribution**
- For a well-behaved chain, the distribution P_t of $\theta^{(t)}$ approaches a stationary distribution π
- Total variation distance is

$$d_{\text{TV}}(P_t, \pi) = \sup_A |P_t(A) - \pi(A)|.$$

It is the largest difference between the probabilities that P_t and π assign to the same event A

- Convergence in total variation means $d_{\text{TV}}(P_t, \pi) \rightarrow 0$
- **Goal in MCMC**: construct T with $\pi(\theta) = p(\theta \mid y)$

- Monte Carlo recap
- **Markov Chain Monte Carlo (MCMC)**
 - Gibbs sampling
 - Metropolis-Hastings
- MCMC Diagnostics
 - Warm-up
 - Convergence
 - S_{eff} , MCSE, and autocorrelation
- Difficult geometries
- Additional material



A stationary distribution: two-state example

Consider a chain with states 1 and 2 and transition matrix

$$T = \begin{pmatrix} 0.8 & 0.2 \\ 0.4 & 0.6 \end{pmatrix}.$$

Each row gives the probabilities of the next state.

- Monte Carlo recap
- Markov Chain Monte Carlo (MCMC)
 - Gibbs sampling
 - Metropolis-Hastings
- MCMC Diagnostics
 - Warm-up
 - Convergence
 - S_{eff} , MCSE, and autocorrelation
- Difficult geometries
- Additional material



A stationary distribution: two-state example

Consider a chain with states 1 and 2 and transition matrix

$$T = \begin{pmatrix} 0.8 & 0.2 \\ 0.4 & 0.6 \end{pmatrix}.$$

Each row gives the probabilities of the next state.

- The distribution $\pi = (2/3, 1/3)$ is stationary because

$$\pi T = (2/3, 1/3) = \pi.$$

- Monte Carlo recap
- Markov Chain Monte Carlo (MCMC)
 - Gibbs sampling
 - Metropolis-Hastings
- MCMC Diagnostics
 - Warm-up
 - Convergence
 - S_{eff} , MCSE, and autocorrelation
- Difficult geometries
- Additional material



A stationary distribution: two-state example

Consider a chain with states 1 and 2 and transition matrix

$$T = \begin{pmatrix} 0.8 & 0.2 \\ 0.4 & 0.6 \end{pmatrix}.$$

Each row gives the probabilities of the next state.

- The distribution $\pi = (2/3, 1/3)$ is stationary because

$$\pi T = (2/3, 1/3) = \pi.$$

- Starting from state 1 gives

$$P_0 = (1, 0), \quad P_1 = (0.8, 0.2), \quad P_2 = (0.72, 0.28), \quad \dots$$

and P_t approaches π

- Monte Carlo recap
- Markov Chain Monte Carlo (MCMC)
 - Gibbs sampling
 - Metropolis-Hastings
- MCMC Diagnostics
 - Warm-up
 - Convergence
 - S_{eff} , MCSE, and autocorrelation
- Difficult geometries
- Additional material



A stationary distribution: two-state example

Consider a chain with states 1 and 2 and transition matrix

$$T = \begin{pmatrix} 0.8 & 0.2 \\ 0.4 & 0.6 \end{pmatrix}.$$

Each row gives the probabilities of the next state.

- The distribution $\pi = (2/3, 1/3)$ is stationary because

$$\pi T = (2/3, 1/3) = \pi.$$

- Starting from state 1 gives

$$P_0 = (1, 0), \quad P_1 = (0.8, 0.2), \quad P_2 = (0.72, 0.28), \quad \dots$$

and P_t approaches π

- Stationary means that applying the transition leaves the distribution unchanged; convergence means approaching it from another starting distribution

- Monte Carlo recap
- Markov Chain Monte Carlo (MCMC)
 - Gibbs sampling
 - Metropolis-Hastings
- MCMC Diagnostics
 - Warm-up
 - Convergence
 - S_{eff} , MCSE, and autocorrelation
- Difficult geometries
- Additional material



UPPSALA
UNIVERSITET

Markov chain Monte Carlo (MCMC)

- Monte Carlo recap
- Markov Chain Monte Carlo (MCMC)
 - Gibbs sampling
 - Metropolis-Hastings
- MCMC Diagnostics
 - Warm-up
 - Convergence
 - S_{eff} , MCSE, and autocorrelation
- Difficult geometries
- Additional material
- Produce draws $\theta^{(t)}$ given $\theta^{(t-1)}$ from a Markov chain, with **stationary distribution** $p(\theta|y)$



Markov chain Monte Carlo (MCMC)

- Monte Carlo recap
- Markov Chain Monte Carlo (MCMC)
 - Gibbs sampling
 - Metropolis-Hastings
- MCMC Diagnostics
 - Warm-up
 - Convergence
 - S_{eff} , MCSE, and autocorrelation
- Difficult geometries
- Additional material
- Produce draws $\theta^{(t)}$ given $\theta^{(t-1)}$ from a Markov chain, with **stationary distribution** $p(\theta|y)$
 - + generic
 - + combine sequence of easier Monte Carlo draws to form a Markov chain



Markov chain Monte Carlo (MCMC)

- Monte Carlo recap
- Markov Chain Monte Carlo (MCMC)
 - Gibbs sampling
 - Metropolis-Hastings
- MCMC Diagnostics
 - Warm-up
 - Convergence
 - S_{eff} , MCSE, and autocorrelation
- Difficult geometries
- Additional material
- Produce draws $\theta^{(t)}$ given $\theta^{(t-1)}$ from a Markov chain, with **stationary distribution** $p(\theta|y)$
 - + generic
 - + combine sequence of easier Monte Carlo draws to form a Markov chain
 - + under suitable irreducibility and aperiodicity conditions, converges in total variation to $p(\theta|y)$ as $t \rightarrow \infty$



Markov chain Monte Carlo (MCMC)

- Monte Carlo recap
- Markov Chain Monte Carlo (MCMC)
 - Gibbs sampling
 - Metropolis-Hastings
- MCMC Diagnostics
 - Warm-up
 - Convergence
 - S_{eff} , MCSE, and autocorrelation
- Difficult geometries
- Additional material

- Produce draws $\theta^{(t)}$ given $\theta^{(t-1)}$ from a Markov chain, with **stationary distribution** $p(\theta|y)$
 - + generic
 - + combine sequence of easier Monte Carlo draws to form a Markov chain
 - + under suitable irreducibility and aperiodicity conditions, converges in total variation to $p(\theta|y)$ as $t \rightarrow \infty$
 - + for well-behaved chains, a central limit theorem holds for Monte Carlo averages



Markov chain Monte Carlo (MCMC)

- Monte Carlo recap
- Markov Chain Monte Carlo (MCMC)
 - Gibbs sampling
 - Metropolis-Hastings
- MCMC Diagnostics
 - Warm-up
 - Convergence
 - S_{eff} , MCSE, and autocorrelation
- Difficult geometries
- Additional material

- Produce draws $\theta^{(t)}$ given $\theta^{(t-1)}$ from a Markov chain, with **stationary distribution** $p(\theta|y)$
 - + generic
 - + combine sequence of easier Monte Carlo draws to form a Markov chain
 - + under suitable irreducibility and aperiodicity conditions, converges in total variation to $p(\theta|y)$ as $t \rightarrow \infty$
 - + for well-behaved chains, a central limit theorem holds for Monte Carlo averages
 - draws are dependent
 - constructing an **efficient** Markov chain is not always easy



- Monte Carlo recap
- **Markov Chain Monte Carlo (MCMC)**
 - Gibbs sampling
 - Metropolis-Hastings
- MCMC Diagnostics
 - Warm-up
 - Convergence
 - S_{eff} , MCSE, and autocorrelation
- Difficult geometries
- Additional material

- Random variables $\theta^{(1)}, \theta^{(2)}, \dots$ where $\theta^{(t)}$ depends only on the previous state $\theta^{(t-1)}$

$$p(\theta^{(t)} | \theta^{(1)}, \dots, \theta^{(t-1)}) = p(\theta^{(t)} | \theta^{(t-1)})$$



- Monte Carlo recap
- **Markov Chain Monte Carlo (MCMC)**
 - Gibbs sampling
 - Metropolis-Hastings
- MCMC Diagnostics
 - Warm-up
 - Convergence
 - S_{eff} , MCSE, and autocorrelation
- Difficult geometries
- Additional material

- Random variables $\theta^{(1)}, \theta^{(2)}, \dots$ where $\theta^{(t)}$ depends only on the previous state $\theta^{(t-1)}$

$$p(\theta^{(t)} | \theta^{(1)}, \dots, \theta^{(t-1)}) = p(\theta^{(t)} | \theta^{(t-1)})$$

- Chain has to be initialized with some starting point $\theta^{(0)}$



- Monte Carlo recap
- **Markov Chain Monte Carlo (MCMC)**
 - Gibbs sampling
 - Metropolis-Hastings
- MCMC Diagnostics
 - Warm-up
 - Convergence
 - S_{eff} , MCSE, and autocorrelation
- Difficult geometries
- Additional material

- Random variables $\theta^{(1)}, \theta^{(2)}, \dots$ where $\theta^{(t)}$ depends only on the previous state $\theta^{(t-1)}$

$$p(\theta^{(t)} | \theta^{(1)}, \dots, \theta^{(t-1)}) = p(\theta^{(t)} | \theta^{(t-1)})$$

- Chain has to be initialized with some starting point $\theta^{(0)}$
- Choose a transition distribution so the **stationary distribution** of the Markov chain is $p(\theta|y)$



UPPSALA UNIVERSITET

- Monte Carlo recap
- Markov Chain Monte Carlo (MCMC)
 - Gibbs sampling
 - Metropolis-Hastings
- MCMC Diagnostics
 - Warm-up
 - Convergence
 - S_{eff} , MCSE, and autocorrelation
- Difficult geometries
- Additional material

Subsection 1

Gibbs sampling



- Alternate sampling from conditional distributions
- For two parameters, one complete sweep is

$$\theta_1^{(t)} \sim p(\theta_1 \mid \theta_2^{(t-1)}, y),$$

$$\theta_2^{(t)} \sim p(\theta_2 \mid \theta_1^{(t)}, y).$$

- Monte Carlo recap
- Markov Chain Monte Carlo (MCMC)
 - Gibbs sampling
 - Metropolis-Hastings
- MCMC Diagnostics
 - Warm-up
 - Convergence
 - S_{eff} , MCSE, and autocorrelation
- Difficult geometries
- Additional material



- Alternate sampling from conditional distributions
- For two parameters, one complete sweep is

$$\theta_1^{(t)} \sim p(\theta_1 \mid \theta_2^{(t-1)}, y),$$

$$\theta_2^{(t)} \sim p(\theta_2 \mid \theta_1^{(t)}, y).$$

- More generally, update $j = 1, \dots, J$ in order using

$$\theta_j^{(t)} \sim p\left(\theta_j \mid \theta_1^{(t)}, \dots, \theta_{j-1}^{(t)}, \theta_{j+1}^{(t-1)}, \dots, \theta_J^{(t-1)}, y\right).$$

- Monte Carlo recap
- Markov Chain Monte Carlo (MCMC)
 - Gibbs sampling
 - Metropolis-Hastings
- MCMC Diagnostics
 - Warm-up
 - Convergence
 - S_{eff} , MCSE, and autocorrelation
- Difficult geometries
- Additional material



- Alternate sampling from conditional distributions
- For two parameters, one complete sweep is

$$\theta_1^{(t)} \sim p(\theta_1 \mid \theta_2^{(t-1)}, y),$$

$$\theta_2^{(t)} \sim p(\theta_2 \mid \theta_1^{(t)}, y).$$

- More generally, update $j = 1, \dots, J$ in order using

$$\theta_j^{(t)} \sim p\left(\theta_j \mid \theta_1^{(t)}, \dots, \theta_{j-1}^{(t)}, \theta_{j+1}^{(t-1)}, \dots, \theta_J^{(t-1)}, y\right).$$

- Under standard conditions, the distribution of $\theta^{(t)}$ converges to $p(\theta \mid y)$ as $t \rightarrow \infty$
- Several parameters can be updated together in a block

- Monte Carlo recap
- Markov Chain Monte Carlo (MCMC)
 - Gibbs sampling
 - Metropolis-Hastings
- MCMC Diagnostics
 - Warm-up
 - Convergence
 - S_{eff} , MCSE, and autocorrelation
- Difficult geometries
- Additional material



Example: Bivariate Normal using Gibbs

For n independent observations $y_i \sim \mathcal{N}(\theta, \Sigma)$ and a flat prior, let $\bar{y}$ be the sample mean. Then

$$\theta \mid y \sim \mathcal{N}\left(\bar{y}, \frac{1}{n}\Sigma\right), \quad \Sigma = \begin{pmatrix} \sigma_1^2 & \rho\sigma_1\sigma_2 \\ \rho\sigma_1\sigma_2 & \sigma_2^2 \end{pmatrix}.$$

The known Σ has observation standard deviations σ_1, σ_2 and correlation ρ . The posterior standard deviations are $\sigma_j/\sqrt{n}$.

- Monte Carlo recap
- Markov Chain Monte Carlo (MCMC)
 - Gibbs sampling
 - Metropolis-Hastings
- MCMC Diagnostics
 - Warm-up
 - Convergence
 - S_{eff} , MCSE, and autocorrelation
- Difficult geometries
- Additional material



Example: Bivariate Normal using Gibbs

For n independent observations $y_i \sim \mathcal{N}(\theta, \Sigma)$ and a flat prior, let $\bar{y}$ be the sample mean. Then

$$\theta \mid y \sim \mathcal{N}\left(\bar{y}, \frac{1}{n}\Sigma\right), \quad \Sigma = \begin{pmatrix} \sigma_1^2 & \rho\sigma_1\sigma_2 \\ \rho\sigma_1\sigma_2 & \sigma_2^2 \end{pmatrix}.$$

The known Σ has observation standard deviations σ_1, σ_2 and correlation ρ . The posterior standard deviations are $\sigma_j/\sqrt{n}$.

One **Gibbs sweep**, given $(\theta_1^{(t-1)}, \theta_2^{(t-1)})$:

$$\theta_1^{(t)} \sim \mathcal{N}\left(\bar{y}_1 + \rho \frac{\sigma_1}{\sigma_2} (\theta_2^{(t-1)} - \bar{y}_2), \frac{\sigma_1^2}{n} (1 - \rho^2)\right),$$

$$\theta_2^{(t)} \sim \mathcal{N}\left(\bar{y}_2 + \rho \frac{\sigma_2}{\sigma_1} (\theta_1^{(t)} - \bar{y}_1), \frac{\sigma_2^2}{n} (1 - \rho^2)\right).$$


- Monte Carlo recap
- Markov Chain Monte Carlo (MCMC)
 - Gibbs sampling
 - Metropolis-Hastings
- MCMC Diagnostics
 - Warm-up
 - Convergence
 - S_{eff} , MCSE, and autocorrelation
- Difficult geometries
- Additional material



UPPSALA
UNIVERSITET

- Monte Carlo recap
- Markov Chain Monte Carlo (MCMC)
 - Gibbs sampling
 - Metropolis-Hastings
- MCMC Diagnostics
 - Warm-up
 - Convergence
 - S_{eff} , MCSE, and autocorrelation
- Difficult geometries
- Additional material

Gibbs sampling



demo



UPPSALA
UNIVERSITET

Gibbs sampling

- Monte Carlo recap
 - Markov Chain Monte Carlo (MCMC)
 - Gibbs sampling
 - Metropolis-Hastings
 - MCMC Diagnostics
 - Warm-up
 - Convergence
 - S_{eff} , MCSE, and autocorrelation
 - Difficult geometries
 - Additional material
- When the full conditional distributions have standard forms, exact Gibbs updates are easy to draw



- Monte Carlo recap
 - Markov Chain Monte Carlo (MCMC)
 - Gibbs sampling
 - Metropolis-Hastings
 - MCMC Diagnostics
 - Warm-up
 - Convergence
 - S_{eff} , MCSE, and autocorrelation
 - Difficult geometries
 - Additional material
- When the full conditional distributions have standard forms, exact Gibbs updates are easy to draw
 - Exact Gibbs updates have no proposal scale or rejection step
 - blocking, parameterization, and update order can still strongly affect efficiency



- Monte Carlo recap
 - Markov Chain Monte Carlo (MCMC)
 - Gibbs sampling
 - Metropolis-Hastings
 - MCMC Diagnostics
 - Warm-up
 - Convergence
 - S_{eff} , MCSE, and autocorrelation
 - Difficult geometries
 - Additional material
- When the full conditional distributions have standard forms, exact Gibbs updates are easy to draw
 - Exact Gibbs updates have no proposal scale or rejection step
 - blocking, parameterization, and update order can still strongly affect efficiency
 - For a nonstandard full conditional, use a Metropolis–Hastings update within the Gibbs sweep



- Monte Carlo recap
- Markov Chain Monte Carlo (MCMC)
 - Gibbs sampling
 - Metropolis-Hastings
- MCMC Diagnostics
 - Warm-up
 - Convergence
 - S_{eff} , MCSE, and autocorrelation
- Difficult geometries
- Additional material

- When the full conditional distributions have standard forms, exact Gibbs updates are easy to draw
- Exact Gibbs updates have no proposal scale or rejection step
 - blocking, parameterization, and update order can still strongly affect efficiency
- For a nonstandard full conditional, use a Metropolis–Hastings update within the Gibbs sweep
- Several parameters can be updated in blocks (*blocking*)




- Monte Carlo recap
- Markov Chain Monte Carlo (MCMC)
 - Gibbs sampling
 - Metropolis-Hastings
- MCMC Diagnostics
 - Warm-up
 - Convergence
 - S_{eff} , MCSE, and autocorrelation
- Difficult geometries
- Additional material

- When the full conditional distributions have standard forms, exact Gibbs updates are easy to draw
- Exact Gibbs updates have no proposal scale or rejection step
 - blocking, parameterization, and update order can still strongly affect efficiency
- For a nonstandard full conditional, use a Metropolis–Hastings update within the Gibbs sweep
- Several parameters can be updated in blocks (*blocking*)
- Slow if parameters are highly dependent in the posterior...



- Monte Carlo recap
- Markov Chain Monte Carlo (MCMC)
 - Gibbs sampling
 - Metropolis-Hastings
- MCMC Diagnostics
 - Warm-up
 - Convergence
 - S_{eff} , MCSE, and autocorrelation
- Difficult geometries
- Additional material



demo



UPPSALA UNIVERSITET

Pair discussion: correlated Gibbs sampling

- Monte Carlo recap
- Markov Chain Monte Carlo (MCMC)
 - Gibbs sampling
 - Metropolis-Hastings
- MCMC Diagnostics
 - Warm-up
 - Convergence
 - S_{eff} , MCSE, and autocorrelation
- Difficult geometries
- Additional material

- Both Gibbs samplers accept every update. Why does the sampler for the correlated target nevertheless produce fewer effective draws?
- What would you expect its trace plot and autocorrelation function to look like?



UPPSALA UNIVERSITET

- Monte Carlo recap
- Markov Chain Monte Carlo (MCMC)
 - Gibbs sampling
 - Metropolis-Hastings
- MCMC Diagnostics
 - Warm-up
 - Convergence
 - S_{eff} , MCSE, and autocorrelation
- Difficult geometries
- Additional material

Pair discussion: correlated Gibbs sampling

- Both Gibbs samplers accept every update. Why does the sampler for the correlated target nevertheless produce fewer effective draws?
- What would you expect its trace plot and autocorrelation function to look like?

After discussion:

- Each coordinate update moves only horizontally or vertically, while the posterior mass lies along a narrow diagonal region.
- The chain therefore moves slowly along that region, producing high autocorrelation and a small effective sample size.



Sampling conditional vs joint

- Monte Carlo recap
 - Markov Chain Monte Carlo (MCMC)
 - Gibbs sampling
 - Metropolis-Hastings
 - MCMC Diagnostics
 - Warm-up
 - Convergence
 - S_{eff} , MCSE, and autocorrelation
 - Difficult geometries
 - Additional material
- How about sampling θ jointly?
 - e.g. it is easy to sample from multivariate normal



Sampling conditional vs joint

- Monte Carlo recap
 - Markov Chain Monte Carlo (MCMC)
 - Gibbs sampling
 - Metropolis-Hastings
 - MCMC Diagnostics
 - Warm-up
 - Convergence
 - S_{eff} , MCSE, and autocorrelation
 - Difficult geometries
 - Additional material
- How about sampling θ jointly?
 - e.g. it is easy to sample from multivariate normal
 - If direct joint sampling is available, use it
 - the draws are independent, so a Markov chain is unnecessary



Sampling conditional vs joint

- Monte Carlo recap
 - Markov Chain Monte Carlo (MCMC)
 - Gibbs sampling
 - Metropolis-Hastings
 - MCMC Diagnostics
 - Warm-up
 - Convergence
 - S_{eff} , MCSE, and autocorrelation
 - Difficult geometries
 - Additional material
- How about sampling θ jointly?
 - e.g. it is easy to sample from multivariate normal
 - If direct joint sampling is available, use it
 - the draws are independent, so a Markov chain is unnecessary
 - MCMC is useful when direct joint sampling is unavailable but full conditionals or local proposals are manageable



UPPSALA
UNIVERSITET

- Monte Carlo recap
- Markov Chain Monte Carlo (MCMC)
 - Gibbs sampling
 - Metropolis-Hastings
- MCMC Diagnostics
 - Warm-up
 - Convergence
 - S_{eff} , MCSE, and autocorrelation
- Difficult geometries
- Additional material

Subsection 2

Metropolis-Hastings



- Algorithm

1. choose $\theta^{(0)}$ with $0 < q(\theta^{(0)}) < \infty$

2. $t = 1, 2, \dots$

- (a) pick a proposal θ^* from a **proposal distribution** $J(\theta^*|\theta^{(t-1)})$.

(Proposal distribution has to be symmetric, i.e.

$J(\theta_a|\theta_b) = J(\theta_b|\theta_a)$, for all θ_a, θ_b)

- Monte Carlo recap
- Markov Chain Monte Carlo (MCMC)
 - Gibbs sampling
 - **Metropolis-Hastings**
- MCMC Diagnostics
 - Warm-up
 - Convergence
 - S_{eff} , MCSE, and autocorrelation
- Difficult geometries
- Additional material



- Monte Carlo recap
- Markov Chain Monte Carlo (MCMC)
 - Gibbs sampling
 - Metropolis-Hastings
- MCMC Diagnostics
 - Warm-up
 - Convergence
 - S_{eff} , MCSE, and autocorrelation
- Difficult geometries
- Additional material

The Metropolis algorithm

- Algorithm

1. choose $\theta^{(0)}$ with $0 < q(\theta^{(0)}) < \infty$
2. $t = 1, 2, \dots$
 - (a) pick a proposal θ^* from a **proposal distribution** $J(\theta^*|\theta^{(t-1)})$.
(Proposal distribution has to be symmetric, i.e. $J(\theta_a|\theta_b) = J(\theta_b|\theta_a)$, for all θ_a, θ_b)
 - (b) calculate acceptance ratio

$$r = \frac{p(\theta^*|y)}{p(\theta^{(t-1)}|y)}$$



- Monte Carlo recap
- Markov Chain Monte Carlo (MCMC)
 - Gibbs sampling
 - Metropolis-Hastings
- MCMC Diagnostics
 - Warm-up
 - Convergence
 - S_{eff} , MCSE, and autocorrelation
- Difficult geometries
- Additional material

The Metropolis algorithm

- Algorithm

1. choose $\theta^{(0)}$ with $0 < q(\theta^{(0)}) < \infty$
2. $t = 1, 2, \dots$
 - (a) pick a proposal θ^* from a **proposal distribution** $J(\theta^*|\theta^{(t-1)})$.
(Proposal distribution has to be symmetric, i.e. $J(\theta_a|\theta_b) = J(\theta_b|\theta_a)$, for all θ_a, θ_b)
 - (b) calculate acceptance ratio

$$r = \frac{p(\theta^*|y)}{p(\theta^{(t-1)}|y)}$$

- (c) set

$$\theta^{(t)} = \begin{cases} \theta^* & \text{with probability } \min(r, 1) \\ \theta^{(t-1)} & \text{otherwise} \end{cases}$$



- Monte Carlo recap
- Markov Chain Monte Carlo (MCMC)
 - Gibbs sampling
 - Metropolis-Hastings
- MCMC Diagnostics
 - Warm-up
 - Convergence
 - S_{eff} , MCSE, and autocorrelation
- Difficult geometries
- Additional material

The Metropolis algorithm

- Algorithm

1. choose $\theta^{(0)}$ with $0 < q(\theta^{(0)}) < \infty$
2. $t = 1, 2, \dots$
 - (a) pick a proposal θ^* from a **proposal distribution** $J(\theta^*|\theta^{(t-1)})$.
(Proposal distribution has to be symmetric, i.e. $J(\theta_a|\theta_b) = J(\theta_b|\theta_a)$, for all θ_a, θ_b)
 - (b) calculate acceptance ratio

$$r = \frac{p(\theta^*|y)}{p(\theta^{(t-1)}|y)}$$

- (c) set

$$\theta^{(t)} = \begin{cases} \theta^* & \text{with probability } \min(r, 1) \\ \theta^{(t-1)} & \text{otherwise} \end{cases}$$

i.e., if $p(\theta^*|y) > p(\theta^{(t-1)}|y)$, always accept the proposal and otherwise accept the proposal with probability r



- Monte Carlo recap
- Markov Chain Monte Carlo (MCMC)
 - Gibbs sampling
 - Metropolis-Hastings
- MCMC Diagnostics
 - Warm-up
 - Convergence
 - S_{eff} , MCSE, and autocorrelation
- Difficult geometries
- Additional material

The Metropolis algorithm

- Algorithm

1. choose $\theta^{(0)}$ with $0 < q(\theta^{(0)}) < \infty$
2. $t = 1, 2, \dots$
 - (a) pick a proposal θ^* from a **proposal distribution** $J(\theta^*|\theta^{(t-1)})$.
(Proposal distribution has to be symmetric, i.e. $J(\theta_a|\theta_b) = J(\theta_b|\theta_a)$, for all θ_a, θ_b)
 - (b) calculate acceptance ratio

(c) set

$$r = \frac{p(\theta^*|y)}{p(\theta^{(t-1)}|y)}$$
$$\theta^{(t)} = \begin{cases} \theta^* & \text{with probability } \min(r, 1) \\ \theta^{(t-1)} & \text{otherwise} \end{cases}$$

- rejection of a proposal increments the time t also by one i.e., the new state is the same as the previous state



- Monte Carlo recap
- Markov Chain Monte Carlo (MCMC)
 - Gibbs sampling
 - Metropolis-Hastings
- MCMC Diagnostics
 - Warm-up
 - Convergence
 - S_{eff} , MCSE, and autocorrelation
- Difficult geometries
- Additional material

The Metropolis algorithm

- Algorithm

1. choose $\theta^{(0)}$ with $0 < q(\theta^{(0)}) < \infty$
2. $t = 1, 2, \dots$
 - (a) pick a proposal θ^* from a **proposal distribution** $J(\theta^*|\theta^{(t-1)})$.
(Proposal distribution has to be symmetric, i.e. $J(\theta_a|\theta_b) = J(\theta_b|\theta_a)$, for all θ_a, θ_b)
 - (b) calculate acceptance ratio

(c) set

$$r = \frac{p(\theta^*|y)}{p(\theta^{(t-1)}|y)}$$
$$\theta^{(t)} = \begin{cases} \theta^* & \text{with probability } \min(r, 1) \\ \theta^{(t-1)} & \text{otherwise} \end{cases}$$

- rejection of a proposal increments the time t also by one i.e., the new state is the same as the previous state
- draw $U \sim \mathcal{U}(0, 1)$ and accept if $U < \min(1, r)$



- Monte Carlo recap
- Markov Chain Monte Carlo (MCMC)
 - Gibbs sampling
 - Metropolis-Hastings
- MCMC Diagnostics
 - Warm-up
 - Convergence
 - S_{eff} , MCSE, and autocorrelation
- Difficult geometries
- Additional material

The Metropolis algorithm

- Algorithm

1. choose $\theta^{(0)}$ with $0 < q(\theta^{(0)}) < \infty$
2. $t = 1, 2, \dots$
 - (a) pick a proposal θ^* from a **proposal distribution** $J(\theta^*|\theta^{(t-1)})$.
(Proposal distribution has to be symmetric, i.e. $J(\theta_a|\theta_b) = J(\theta_b|\theta_a)$, for all θ_a, θ_b)
 - (b) calculate acceptance ratio

(c) set

$$r = \frac{p(\theta^*|y)}{p(\theta^{(t-1)}|y)}$$
$$\theta^{(t)} = \begin{cases} \theta^* & \text{with probability } \min(r, 1) \\ \theta^{(t-1)} & \text{otherwise} \end{cases}$$

- rejection of a proposal increments the time t also by one i.e., the new state is the same as the previous state
- draw $U \sim \mathcal{U}(0, 1)$ and accept if $U < \min(1, r)$
- $p(\theta^*|y)$ and $p(\theta^{(t-1)}|y)$ have the same normalizing constant, so the unnormalized density $q(\cdot)$ can be used:
the constant cancels



Metropolis on the log scale

- For a symmetric proposal, use the unnormalized log posterior

$$\ell(\theta) = \log q(\theta) = \log p(y \mid \theta) + \log p(\theta).$$

- Monte Carlo recap
- Markov Chain Monte Carlo (MCMC)
 - Gibbs sampling
 - Metropolis-Hastings
- MCMC Diagnostics
 - Warm-up
 - Convergence
 - S_{eff} , MCSE, and autocorrelation
- Difficult geometries
- Additional material



Metropolis on the log scale

- For a symmetric proposal, use the unnormalized log posterior

$$\ell(\theta) = \log q(\theta) = \log p(y \mid \theta) + \log p(\theta).$$

- Compute the log density ratio

$$\Delta = \ell(\theta^*) - \ell(\theta^{(t-1)}).$$

- Monte Carlo recap
- Markov Chain Monte Carlo (MCMC)
 - Gibbs sampling
 - Metropolis-Hastings
- MCMC Diagnostics
 - Warm-up
 - Convergence
 - S_{eff} , MCSE, and autocorrelation
- Difficult geometries
- Additional material



Metropolis on the log scale

- For a symmetric proposal, use the unnormalized log posterior

$$\ell(\theta) = \log q(\theta) = \log p(y \mid \theta) + \log p(\theta).$$

- Compute the log density ratio

$$\Delta = \ell(\theta^*) - \ell(\theta^{(t-1)}).$$

- Draw $U \sim \mathcal{U}(0, 1)$ and accept the proposal if

$$\log U < \min(0, \Delta).$$

Otherwise repeat the current state: $\theta^{(t)} = \theta^{(t-1)}$

- Monte Carlo recap
- Markov Chain Monte Carlo (MCMC)
 - Gibbs sampling
 - Metropolis-Hastings
- MCMC Diagnostics
 - Warm-up
 - Convergence
 - S_{eff} , MCSE, and autocorrelation
- Difficult geometries
- Additional material



Metropolis on the log scale

- For a symmetric proposal, use the unnormalized log posterior

$$\ell(\theta) = \log q(\theta) = \log p(y \mid \theta) + \log p(\theta).$$

- Compute the log density ratio

$$\Delta = \ell(\theta^*) - \ell(\theta^{(t-1)}).$$

- Draw $U \sim \mathcal{U}(0, 1)$ and accept the proposal if

$$\log U < \min(0, \Delta).$$

Otherwise repeat the current state: $\theta^{(t)} = \theta^{(t-1)}$

- This avoids evaluating tiny densities or exponentiating large log density ratios

- Monte Carlo recap
- Markov Chain Monte Carlo (MCMC)
 - Gibbs sampling
 - Metropolis-Hastings
- MCMC Diagnostics
 - Warm-up
 - Convergence
 - S_{eff} , MCSE, and autocorrelation
- Difficult geometries
- Additional material



Rejection sampling vs. Metropolis rejection

- In rejection sampling, a rejected proposal is discarded
 - only accepted proposals are included in the output
 - the next proposal is drawn independently from g

- Monte Carlo recap
- Markov Chain Monte Carlo (MCMC)
 - Gibbs sampling
 - Metropolis-Hastings
- MCMC Diagnostics
 - Warm-up
 - Convergence
 - S_{eff} , MCSE, and autocorrelation
- Difficult geometries
- Additional material



Rejection sampling vs. Metropolis rejection

- In rejection sampling, a rejected proposal is discarded
 - only accepted proposals are included in the output
 - the next proposal is drawn independently from g
- In the Metropolis algorithm, a rejected proposal is discarded, but the current state is repeated

$$\theta^{(t)} = \theta^{(t-1)}.$$

- Monte Carlo recap
- Markov Chain Monte Carlo (MCMC)
 - Gibbs sampling
 - Metropolis-Hastings
- MCMC Diagnostics
 - Warm-up
 - Convergence
 - S_{eff} , MCSE, and autocorrelation
- Difficult geometries
- Additional material

- the repeated state is a valid draw from the Markov chain
- it must be retained when computing posterior summaries



Rejection sampling vs. Metropolis rejection

- In rejection sampling, a rejected proposal is discarded
 - only accepted proposals are included in the output
 - the next proposal is drawn independently from g
- In the Metropolis algorithm, a rejected proposal is discarded, but the current state is repeated

$$\theta^{(t)} = \theta^{(t-1)}.$$

- the repeated state is a valid draw from the Markov chain
 - it must be retained when computing posterior summaries
- Keeping only accepted Metropolis proposals changes the distribution and generally gives incorrect estimates

- Monte Carlo recap
- Markov Chain Monte Carlo (MCMC)
 - Gibbs sampling
 - Metropolis-Hastings
- MCMC Diagnostics
 - Warm-up
 - Convergence
 - S_{eff} , MCSE, and autocorrelation
- Difficult geometries
- Additional material



Consider a random-walk Metropolis proposal

$$\theta^* \sim \mathcal{N}(\theta^{(t-1)}, \sigma_p^2 I).$$

- Predict what happens when σ_p is much too small.
- Predict what happens when σ_p is much too large.
- For each case, what do you expect for the acceptance rate, trace plot, and autocorrelation?

- Monte Carlo recap
- Markov Chain Monte Carlo (MCMC)
 - Gibbs sampling
 - Metropolis-Hastings
- MCMC Diagnostics
 - Warm-up
 - Convergence
 - S_{eff} , MCSE, and autocorrelation
- Difficult geometries
- Additional material



Proposal scale: discussion

Consider a random-walk Metropolis proposal

$$\theta^* \sim \mathcal{N}(\theta^{(t-1)}, \sigma_p^2 I).$$

- Predict what happens when σ_p is much too small.
- Predict what happens when σ_p is much too large.
- For each case, what do you expect for the acceptance rate, trace plot, and autocorrelation?

After discussion: Small proposals are often accepted but move slowly. Large proposals move farther when accepted, but are often rejected. Both choices can give high autocorrelation and low ESS.

- Monte Carlo recap
- Markov Chain Monte Carlo (MCMC)
 - Gibbs sampling
 - Metropolis-Hastings
- MCMC Diagnostics
 - Warm-up
 - Convergence
 - S_{eff} , MCSE, and autocorrelation
- Difficult geometries
- Additional material



- Monte Carlo recap
- Markov Chain Monte Carlo (MCMC)
 - Gibbs sampling
 - Metropolis-Hastings
- MCMC Diagnostics
 - Warm-up
 - Convergence
 - S_{eff} , MCSE, and autocorrelation
- Difficult geometries
- Additional material

- Example: one bivariate observation $y = (y_1, y_2)$
 - $y \mid \theta \sim \mathcal{N}(\theta, \Sigma)$, with known covariance Σ and a flat prior $p(\theta) \propto 1$

$$\begin{pmatrix} \theta_1 \\ \theta_2 \end{pmatrix} \Big| y \sim \mathcal{N} \left(\begin{pmatrix} y_1 \\ y_2 \end{pmatrix}, \begin{pmatrix} 1 & \rho \\ \rho & 1 \end{pmatrix} \right)$$

- proposal distribution $J(\theta^* | \theta^{(t-1)}) = \mathcal{N}(\theta^* | \theta^{(t-1)}, \sigma_p^2 I)$

demo



Why the Metropolis algorithm works

- Intuitively, the chain spends more time in high-density regions
 - proposals to higher density are always accepted
 - proposals to lower density are accepted only sometimes

- Monte Carlo recap
- Markov Chain Monte Carlo (MCMC)
 - Gibbs sampling
 - Metropolis-Hastings
- MCMC Diagnostics
 - Warm-up
 - Convergence
 - S_{eff} , MCSE, and autocorrelation
- Difficult geometries
- Additional material



Why the Metropolis algorithm works

- Monte Carlo recap
- Markov Chain Monte Carlo (MCMC)
 - Gibbs sampling
 - Metropolis-Hastings
- MCMC Diagnostics
 - Warm-up
 - Convergence
 - S_{eff} , MCSE, and autocorrelation
- Difficult geometries
- Additional material

- Intuitively, the chain spends more time in high-density regions
 - proposals to higher density are always accepted
 - proposals to lower density are accepted only sometimes
- Detailed balance will show that $p(\theta | y)$ is a stationary distribution of the Metropolis chain



Why the Metropolis algorithm works

- Monte Carlo recap
- Markov Chain Monte Carlo (MCMC)
 - Gibbs sampling
 - Metropolis-Hastings
- MCMC Diagnostics
 - Warm-up
 - Convergence
 - S_{eff} , MCSE, and autocorrelation
- Difficult geometries
- Additional material

- Intuitively, the chain spends more time in high-density regions
 - proposals to higher density are always accepted
 - proposals to lower density are accepted only sometimes
- Detailed balance will show that $p(\theta | y)$ is a stationary distribution of the Metropolis chain
- Convergence to that stationary distribution also requires a well-behaved chain
 - the proposal must allow the chain to reach all relevant regions of the target distribution
 - the chain must not be trapped in a deterministic cycle



Why the Metropolis algorithm works

- Monte Carlo recap
- Markov Chain Monte Carlo (MCMC)
 - Gibbs sampling
 - Metropolis-Hastings
- MCMC Diagnostics
 - Warm-up
 - Convergence
 - S_{eff} , MCSE, and autocorrelation
- Difficult geometries
- Additional material

- Intuitively, the chain spends more time in high-density regions
 - proposals to higher density are always accepted
 - proposals to lower density are accepted only sometimes
- Detailed balance will show that $p(\theta | y)$ is a stationary distribution of the Metropolis chain
- Convergence to that stationary distribution also requires a well-behaved chain
 - the proposal must allow the chain to reach all relevant regions of the target distribution
 - the chain must not be trapped in a deterministic cycle
- These ideas are formalized by **irreducibility** and **aperiodicity**



Why the Metropolis algorithm works

- Monte Carlo recap
- Markov Chain Monte Carlo (MCMC)
 - Gibbs sampling
 - Metropolis-Hastings
- MCMC Diagnostics
 - Warm-up
 - Convergence
 - S_{eff} , MCSE, and autocorrelation
- Difficult geometries
- Additional material

- Intuitively, the chain spends more time in high-density regions
 - proposals to higher density are always accepted
 - proposals to lower density are accepted only sometimes
- Detailed balance will show that $p(\theta | y)$ is a stationary distribution of the Metropolis chain
- Convergence to that stationary distribution also requires a well-behaved chain
 - the proposal must allow the chain to reach all relevant regions of the target distribution
 - the chain must not be trapped in a deterministic cycle
- These ideas are formalized by **irreducibility** and **aperiodicity**
- They depend on both the target and the proposal; they do not follow merely because the target is a proper distribution



Detailed balance for the Metropolis algorithm

- For a symmetric proposal $J(\theta' | \theta)$, the Metropolis acceptance probability is

$$\alpha(\theta, \theta') = \min\left(1, \frac{p(\theta' | y)}{p(\theta | y)}\right).$$

- Monte Carlo recap
- Markov Chain Monte Carlo (MCMC)
 - Gibbs sampling
 - Metropolis-Hastings
- MCMC Diagnostics
 - Warm-up
 - Convergence
 - S_{eff} , MCSE, and autocorrelation
- Difficult geometries
- Additional material



Detailed balance for the Metropolis algorithm

- For a symmetric proposal $J(\theta' | \theta)$, the Metropolis acceptance probability is

$$\alpha(\theta, \theta') = \min\left(1, \frac{p(\theta' | y)}{p(\theta | y)}\right).$$

- For $\theta' \neq \theta$, the probability flow from θ to θ' is

$$\begin{aligned} p(\theta | y) J(\theta' | \theta) \alpha(\theta, \theta') \\ = J(\theta' | \theta) \min(p(\theta | y), p(\theta' | y)). \end{aligned}$$

- Monte Carlo recap
- Markov Chain Monte Carlo (MCMC)
 - Gibbs sampling
 - Metropolis-Hastings
- MCMC Diagnostics
 - Warm-up
 - Convergence
 - S_{eff} , MCSE, and autocorrelation
- Difficult geometries
- Additional material



Detailed balance for the Metropolis algorithm

- For a symmetric proposal $J(\theta' | \theta)$, the Metropolis acceptance probability is

$$\alpha(\theta, \theta') = \min\left(1, \frac{p(\theta' | y)}{p(\theta | y)}\right).$$

- For $\theta' \neq \theta$, the probability flow from θ to θ' is

$$\begin{aligned} p(\theta | y) J(\theta' | \theta) \alpha(\theta, \theta') \\ = J(\theta' | \theta) \min(p(\theta | y), p(\theta' | y)). \end{aligned}$$

- Symmetry of J makes this expression unchanged when θ and θ' are exchanged. Hence

$$p(\theta | y) T(\theta' | \theta) = p(\theta' | y) T(\theta | \theta').$$

- Monte Carlo recap
- Markov Chain Monte Carlo (MCMC)
 - Gibbs sampling
 - Metropolis-Hastings
- MCMC Diagnostics
 - Warm-up
 - Convergence
 - S_{eff} , MCSE, and autocorrelation
- Difficult geometries
- Additional material



- Monte Carlo recap
- Markov Chain Monte Carlo (MCMC)
 - Gibbs sampling
 - Metropolis-Hastings
- MCMC Diagnostics
 - Warm-up
 - Convergence
 - S_{eff} , MCSE, and autocorrelation
- Difficult geometries
- Additional material

Detailed balance for the Metropolis algorithm

- For a symmetric proposal $J(\theta' | \theta)$, the Metropolis acceptance probability is

$$\alpha(\theta, \theta') = \min\left(1, \frac{p(\theta' | y)}{p(\theta | y)}\right).$$

- For $\theta' \neq \theta$, the probability flow from θ to θ' is

$$\begin{aligned} p(\theta | y) J(\theta' | \theta) \alpha(\theta, \theta') \\ = J(\theta' | \theta) \min(p(\theta | y), p(\theta' | y)). \end{aligned}$$

- Symmetry of J makes this expression unchanged when θ and θ' are exchanged. Hence

$$p(\theta | y) T(\theta' | \theta) = p(\theta' | y) T(\theta | \theta').$$

- The self-transition caused by rejection balances trivially. Thus detailed balance holds for the full transition kernel T , and $p(\theta | y)$ is stationary



- Monte Carlo recap
- Markov Chain Monte Carlo (MCMC)
 - Gibbs sampling
 - Metropolis-Hastings
- MCMC Diagnostics
 - Warm-up
 - Convergence
 - S_{eff} , MCSE, and autocorrelation
- Difficult geometries
- Additional material

- Generalization of Metropolis algorithm for **non-symmetric** proposal distributions
 - acceptance ratio includes ratio of proposal distributions

$$r = \frac{p(\theta^* | y) J(\theta^{(t-1)} | \theta^*)}{p(\theta^{(t-1)} | y) J(\theta^* | \theta^{(t-1)})}$$



- The ideal proposal is the target distribution itself
 - $J(\theta^*|\theta) \equiv p(\theta^*|y)$ gives independent draws with acceptance probability 1, but is not usually feasible

- Monte Carlo recap
- Markov Chain Monte Carlo (MCMC)
 - Gibbs sampling
 - Metropolis-Hastings
- MCMC Diagnostics
 - Warm-up
 - Convergence
 - S_{eff} , MCSE, and autocorrelation
- Difficult geometries
- Additional material



UPPSALA UNIVERSITET

- Monte Carlo recap
- Markov Chain Monte Carlo (MCMC)
 - Gibbs sampling
 - Metropolis-Hastings
- MCMC Diagnostics
 - Warm-up
 - Convergence
 - S_{eff} , MCSE, and autocorrelation
- Difficult geometries
- Additional material

Metropolis-Hastings algorithm

- The ideal proposal is the target distribution itself
 - $J(\theta^*|\theta) \equiv p(\theta^*|y)$ gives independent draws with acceptance probability 1, but is not usually feasible
- A good proposal resembles the target distribution



- Monte Carlo recap
- Markov Chain Monte Carlo (MCMC)
 - Gibbs sampling
 - Metropolis-Hastings
- MCMC Diagnostics
 - Warm-up
 - Convergence
 - S_{eff} , MCSE, and autocorrelation
- Difficult geometries
- Additional material

- The ideal proposal is the target distribution itself
 - $J(\theta^*|\theta) \equiv p(\theta^*|y)$ gives independent draws with acceptance probability 1, but is not usually feasible
- A good proposal resembles the target distribution
- Proposal scale controls movement through the target
 - small scale: many accepted steps, but each step moves little
 - large scale: longer proposed steps, but many are rejected



- Monte Carlo recap
- Markov Chain Monte Carlo (MCMC)
 - Gibbs sampling
 - Metropolis-Hastings
- MCMC Diagnostics
 - Warm-up
 - Convergence
 - S_{eff} , MCSE, and autocorrelation
- Difficult geometries
- Additional material

Metropolis-Hastings algorithm

- The ideal proposal is the target distribution itself
 - $J(\theta^*|\theta) \equiv p(\theta^*|y)$ gives independent draws with acceptance probability 1, but is not usually feasible
- A good proposal resembles the target distribution
- Proposal scale controls movement through the target
 - small scale: many accepted steps, but each step moves little
 - large scale: longer proposed steps, but many are rejected
- There is no universal optimal acceptance rate
 - it depends on the target, dimension, proposal, and quantity of interest
 - compare proposals using ESS per computation time or MCSE



- Monte Carlo recap
- Markov Chain Monte Carlo (MCMC)
 - Gibbs sampling
 - Metropolis-Hastings
- MCMC Diagnostics
 - Warm-up
 - Convergence
 - S_{eff} , MCSE, and autocorrelation
- Difficult geometries
- Additional material

Metropolis-Hastings algorithm

- The ideal proposal is the target distribution itself
 - $J(\theta^*|\theta) \equiv p(\theta^*|y)$ gives independent draws with acceptance probability 1, but is not usually feasible
- A good proposal resembles the target distribution
- Proposal scale controls movement through the target
 - small scale: many accepted steps, but each step moves little
 - large scale: longer proposed steps, but many are rejected
- There is no universal optimal acceptance rate
 - it depends on the target, dimension, proposal, and quantity of interest
 - compare proposals using ESS per computation time or MCSE
- [Interactive demonstration](#)



- Update one component while holding θ_{-j} fixed, proposing

$$\theta_j^* \sim p(\theta_j \mid \theta_{-j}, y).$$

- Monte Carlo recap
- Markov Chain Monte Carlo (MCMC)
 - Gibbs sampling
 - Metropolis-Hastings
- MCMC Diagnostics
 - Warm-up
 - Convergence
 - S_{eff} , MCSE, and autocorrelation
- Difficult geometries
- Additional material



- Update one component while holding θ_{-j} fixed, proposing

$$\theta_j^* \sim p(\theta_j \mid \theta_{-j}, y).$$

- The Metropolis–Hastings acceptance ratio is

$$\begin{aligned} r &= \frac{p(\theta_j^*, \theta_{-j} \mid y) p(\theta_j \mid \theta_{-j}, y)}{p(\theta_j, \theta_{-j} \mid y) p(\theta_j^* \mid \theta_{-j}, y)} \\ &= 1, \end{aligned}$$

because each joint posterior factors into the same marginal for θ_{-j} times its full conditional

- Monte Carlo recap
- Markov Chain Monte Carlo (MCMC)
 - Gibbs sampling
 - Metropolis-Hastings
- MCMC Diagnostics
 - Warm-up
 - Convergence
 - S_{eff} , MCSE, and autocorrelation
- Difficult geometries
- Additional material



- Update one component while holding θ_{-j} fixed, proposing

$$\theta_j^* \sim p(\theta_j \mid \theta_{-j}, y).$$

- The Metropolis–Hastings acceptance ratio is

$$\begin{aligned} r &= \frac{p(\theta_j^*, \theta_{-j} \mid y) p(\theta_j \mid \theta_{-j}, y)}{p(\theta_j, \theta_{-j} \mid y) p(\theta_j^* \mid \theta_{-j}, y)} \\ &= 1, \end{aligned}$$

because each joint posterior factors into the same marginal for θ_{-j} times its full conditional


- Therefore every exact Gibbs proposal is accepted

- Monte Carlo recap
- Markov Chain Monte Carlo (MCMC)
 - Gibbs sampling
 - Metropolis-Hastings
- MCMC Diagnostics
 - Warm-up
 - Convergence
 - S_{eff} , MCSE, and autocorrelation
- Difficult geometries
- Additional material




- Monte Carlo recap
- Markov Chain Monte Carlo (MCMC)
 - Gibbs sampling
 - Metropolis-Hastings
- MCMC Diagnostics
 - Warm-up
 - Convergence
 - S_{eff} , MCSE, and autocorrelation
- Difficult geometries
- Additional material

- Random-walk Metropolis usually does not scale well to high dimensions
 - if the shape does not match the whole distribution, efficiency drops



• Draws — Steps of the sampler — 90% HPI



• Draws — Steps of the sampler — 90% HPI



UPPSALA
UNIVERSITET

- Monte Carlo recap
- Markov Chain Monte Carlo (MCMC)
 - Gibbs sampling
 - Metropolis-Hastings
- MCMC Diagnostics
 - Warm-up
 - Convergence
 - S_{eff} , MCSE, and autocorrelation
- Difficult geometries
- Additional material

Section 3

MCMC Diagnostics



UPPSALA
UNIVERSITET

What we want to know

- Monte Carlo recap
- Markov Chain Monte Carlo (MCMC)
 - Gibbs sampling
 - Metropolis-Hastings
- MCMC Diagnostics
 - Warm-up
 - Convergence
 - S_{eff} , MCSE, and autocorrelation
- Difficult geometries
- Additional material

1. Do the chains show evidence of poor **convergence** or mixing?



What we want to know

- Monte Carlo recap
- Markov Chain Monte Carlo (MCMC)
 - Gibbs sampling
 - Metropolis-Hastings
- MCMC Diagnostics
 - Warm-up
 - Convergence
 - S_{eff} , MCSE, and autocorrelation
- Difficult geometries
- Additional material

1. Do the chains show evidence of poor **convergence** or mixing?
2. How many **effective** samples do I have?



UPPSALA UNIVERSITET

- Monte Carlo recap
- Markov Chain Monte Carlo (MCMC)
 - Gibbs sampling
 - Metropolis-Hastings
- MCMC Diagnostics
 - Warm-up
 - Convergence
 - S_{eff} , MCSE, and autocorrelation
- Difficult geometries
- Additional material

Subsection 1

Warm-up



Warm-up


- Early draws can be strongly influenced by the starting value

- Monte Carlo recap
- Markov Chain Monte Carlo (MCMC)
 - Gibbs sampling
 - Metropolis-Hastings
- MCMC Diagnostics
 - Warm-up
 - Convergence
 - S_{eff} , MCSE, and autocorrelation
- Difficult geometries
- Additional material



Warm-up

- Early draws can be strongly influenced by the starting value
 - the chain may not yet have reached the typical posterior region




- Monte Carlo recap
- Markov Chain Monte Carlo (MCMC)
 - Gibbs sampling
 - Metropolis-Hastings
- MCMC Diagnostics
 - Warm-up
 - Convergence
 - S_{eff} , MCSE, and autocorrelation
- Difficult geometries
- Additional material



- Monte Carlo recap
- Markov Chain Monte Carlo (MCMC)
 - Gibbs sampling
 - Metropolis-Hastings
- MCMC Diagnostics
 - Warm-up
 - Convergence
 - S_{eff} , MCSE, and autocorrelation
- Difficult geometries
- Additional material

Warm-up

- Early draws can be strongly influenced by the starting value
 - the chain may not yet have reached the typical posterior region




- Warm-up is the initial phase before retained sampling
 - tune the proposal during warm-up if needed
 - for our Metropolis samplers, keep the proposal fixed during retained sampling
 - discard warm-up (**burn-in**) draws



Warm-up

- Early draws can be strongly influenced by the starting value
 - the chain may not yet have reached the typical posterior region



- Warm-up is the initial phase before retained sampling
 - tune the proposal during warm-up if needed
 - for our Metropolis samplers, keep the proposal fixed during retained sampling
 - discard warm-up (**burn-in**) draws
- Discarding warm-up does not guarantee good sampling
 - the retained chains can still mix slowly or miss posterior regions
 - convergence and efficiency diagnostics are still required



UPPSALA
UNIVERSITET

- Monte Carlo recap
- Markov Chain Monte Carlo (MCMC)
 - Gibbs sampling
 - Metropolis-Hastings
- MCMC Diagnostics
 - Warm-up
 - **Convergence**
 - S_{eff} , MCSE, and autocorrelation
- Difficult geometries
- Additional material

Subsection 2

Convergence



UPPSALA
UNIVERSITET

Assessing convergence

- Several Markov chains make convergence diagnostics easier


- Monte Carlo recap
- Markov Chain Monte Carlo (MCMC)
 - Gibbs sampling
 - Metropolis-Hastings
- MCMC Diagnostics
 - Warm-up
 - **Convergence**
 - S_{eff} , MCSE, and autocorrelation
- Difficult geometries
- Additional material



- Monte Carlo recap
- Markov Chain Monte Carlo (MCMC)
 - Gibbs sampling
 - Metropolis-Hastings
- MCMC Diagnostics
 - Warm-up
 - **Convergence**
 - S_{eff} , MCSE, and autocorrelation
- Difficult geometries
- Additional material

Assessing convergence


- Several Markov chains make convergence diagnostics easier
- Start chains from different, preferably overdispersed, starting points



- Remove warm-up draws and run chains long enough




Not converged



- Monte Carlo recap
- Markov Chain Monte Carlo (MCMC)
 - Gibbs sampling
 - Metropolis-Hastings
- MCMC Diagnostics
 - Warm-up
 - Convergence
 - S_{eff} , MCSE, and autocorrelation
- Difficult geometries
- Additional material




Visually converged



- Monte Carlo recap
- Markov Chain Monte Carlo (MCMC)
 - Gibbs sampling
 - Metropolis-Hastings
- MCMC Diagnostics
 - Warm-up
 - Convergence
 - S_{eff} , MCSE, and autocorrelation
- Difficult geometries
- Additional material



Visually converged



Visual convergence check is not sufficient

- Monte Carlo recap
- Markov Chain Monte Carlo (MCMC)
 - Gibbs sampling
 - Metropolis-Hastings
- MCMC Diagnostics
 - Warm-up
 - **Convergence**
 - S_{eff} , MCSE, and autocorrelation
- Difficult geometries
- Additional material



$\hat{R}$: comparison of within and between variances of the chains

- $\hat{R}$ or potential scale reduction factor (PSRF)
- Compare means and variances of the chains

- Monte Carlo recap
- Markov Chain Monte Carlo (MCMC)
 - Gibbs sampling
 - Metropolis-Hastings
- MCMC Diagnostics
 - Warm-up
 - Convergence
 - S_{eff} , MCSE, and autocorrelation
- Difficult geometries
- Additional material




$\hat{R}$: comparison of within and between variances of the chains


- Monte Carlo recap
- Markov Chain Monte Carlo (MCMC)
 - Gibbs sampling
 - Metropolis-Hastings
- MCMC Diagnostics
 - Warm-up
 - Convergence
 - S_{eff} , MCSE, and autocorrelation
- Difficult geometries
- Additional material

- $\hat{R}$ or **potential scale reduction factor** (PSRF)
- Compare means and variances of the chains
 W = within-chain variance estimate
 $\widehat{\text{var}}^+$ = total variance estimate

50 warmup, 50 post warmup iterations



Rhat = 1.64






$\hat{R}$: comparison of within and between variances of the chains


- Monte Carlo recap
- Markov Chain Monte Carlo (MCMC)
 - Gibbs sampling
 - Metropolis-Hastings
- MCMC Diagnostics
 - Warm-up
 - Convergence
 - S_{eff} , MCSE, and autocorrelation
- Difficult geometries
- Additional material

- $\hat{R}$ or **potential scale reduction factor** (PSRF)
- Compare means and variances of the chains
 W = within-chain variance estimate
 $\widehat{\text{var}}^+$ = total variance estimate

500 warmup, 500 post warmup iterations



Rhat = 1.08







$\hat{R}$: comparison of within and between variances of the chains

- $\hat{R}$ or **potential scale reduction factor** (PSRF)
- Compare means and variances of the chains
 W = within-chain variance estimate
 $\widehat{\text{var}}^+$ = total variance estimate

5000 warmup, 5000 post warmup iterations



Rhat = 1





Classical $\hat{R}$: interpretation

- The classical diagnostic compares a total variance estimate $\widehat{\text{var}}^+$ with the within-chain variance W

$$\hat{R} = \sqrt{\frac{\widehat{\text{var}}^+}{W}}.$$

- Monte Carlo recap
- Markov Chain Monte Carlo (MCMC)
 - Gibbs sampling
 - Metropolis-Hastings
- MCMC Diagnostics
 - Warm-up
 - **Convergence**
 - S_{eff} , MCSE, and autocorrelation
- Difficult geometries
- Additional material



- The classical diagnostic compares a total variance estimate $\widehat{\text{var}}^+$ with the within-chain variance W

$$\hat{R} = \sqrt{\frac{\widehat{\text{var}}^+}{W}}.$$

- $\hat{R}$ close to 1 means the between-chain and within-chain variation are similar
- $\hat{R} > 1$ indicates that the chains disagree more than expected from their within-chain variation
- Course rule of thumb: aim for $\hat{R} < 1.01$, with values closer to 1 preferred

- Monte Carlo recap
- Markov Chain Monte Carlo (MCMC)
 - Gibbs sampling
 - Metropolis-Hastings
- MCMC Diagnostics
 - Warm-up
 - **Convergence**
 - S_{eff} , MCSE, and autocorrelation
- Difficult geometries
- Additional material



- The classical diagnostic compares a total variance estimate $\widehat{\text{var}}^+$ with the within-chain variance W

$$\hat{R} = \sqrt{\frac{\widehat{\text{var}}^+}{W}}.$$

- $\hat{R}$ close to 1 means the between-chain and within-chain variation are similar
- $\hat{R} > 1$ indicates that the chains disagree more than expected from their within-chain variation
- Course rule of thumb: aim for $\hat{R} < 1.01$, with values closer to 1 preferred
- A value close to 1 is necessary but not sufficient: also inspect trace plots, ESS, and MCSE

- Monte Carlo recap
- Markov Chain Monte Carlo (MCMC)
 - Gibbs sampling
 - Metropolis-Hastings
- MCMC Diagnostics
 - Warm-up
 - **Convergence**
 - S_{eff} , MCSE, and autocorrelation
- Difficult geometries
- Additional material



UPPSALA UNIVERSITET

- Monte Carlo recap
- Markov Chain Monte Carlo (MCMC)
 - Gibbs sampling
 - Metropolis-Hastings
- MCMC Diagnostics
 - Warm-up
 - Convergence
 - S_{eff} , MCSE, and autocorrelation
- Difficult geometries
- Additional material

Subsection 3

S_{eff} , MCSE, and autocorrelation



- Monte Carlo recap
- Markov Chain Monte Carlo (MCMC)
 - Gibbs sampling
 - Metropolis-Hastings
- MCMC Diagnostics
 - Warm-up
 - Convergence
 - S_{eff} , MCSE, and autocorrelation
- Difficult geometries
- Additional material

- Notation: M chains, N retained draws per chain, and $S = MN$ retained draws in total
- For well-behaved chains, Monte Carlo averages converge to posterior expectations

$$E_{p(\theta|y)}[f(\theta)] \approx \frac{1}{S} \sum_{s=1}^S f(\theta^{(s)})$$

- A Markov chain central limit theorem can then be used to estimate Monte Carlo error, but dependence makes this more difficult
 - evaluation of *effective* sample size, S_{eff} .




- Monte Carlo recap
 - Markov Chain Monte Carlo (MCMC)
 - Gibbs sampling
 - Metropolis-Hastings
 - MCMC Diagnostics
 - Warm-up
 - Convergence
 - S_{eff} , MCSE, and autocorrelation
 - Difficult geometries
 - Additional material
- The autocorrelation function measures dependence between draws separated by a given lag
 - high autocorrelation means that successive draws contain less new information
 - it helps compare the efficiency of MCMC methods



- Monte Carlo recap
- Markov Chain Monte Carlo (MCMC)
 - Gibbs sampling
 - Metropolis-Hastings
- MCMC Diagnostics
 - Warm-up
 - Convergence
 - S_{eff} , MCSE, and autocorrelation
- Difficult geometries
- Additional material

Autocorrelation: one worked example

Less correlated target: faster exploration




• Draws — Steps of the sampler — 90% HPI




- Monte Carlo recap
- Markov Chain Monte Carlo (MCMC)
 - Gibbs sampling
 - Metropolis-Hastings
- MCMC Diagnostics
 - Warm-up
 - Convergence
 - S_{eff} , MCSE, and autocorrelation
- Difficult geometries
- Additional material

Autocorrelation: one worked example

Less correlated target: faster exploration



• Draws — Steps of the sampler — 90% HPDI







- Monte Carlo recap
- Markov Chain Monte Carlo (MCMC)
 - Gibbs sampling
 - Metropolis-Hastings
- MCMC Diagnostics
 - Warm-up
 - Convergence
 - S_{eff} , MCSE, and autocorrelation
- Difficult geometries
- Additional material


Autocorrelation: one worked example

Less correlated target: faster exploration



• Draws — Steps of the sampler — 90% HPD



Autocorrelation function






Autocorrelation: comparison example 1

Stronger target correlation, same proposal scale



• Draws — Steps of the sampler — 90% HPD

Autocorrelation function





- Monte Carlo recap
- Markov Chain Monte Carlo (MCMC)
 - Gibbs sampling
 - Metropolis-Hastings
- MCMC Diagnostics
 - Warm-up
 - Convergence
 - S_{eff} , MCSE, and autocorrelation
- Difficult geometries
- Additional material




Autocorrelation: comparison example 2

Same strongly correlated target, larger proposals



• Draws — Steps of the sampler — 90% HPD

Autocorrelation function



- Monte Carlo recap
- Markov Chain Monte Carlo (MCMC)
 - Gibbs sampling
 - Metropolis-Hastings
- MCMC Diagnostics
 - Warm-up
 - Convergence
 - S_{eff} , MCSE, and autocorrelation
- Difficult geometries
- Additional material



From autocorrelation to Monte Carlo error

Consider one stationary chain with $X_s = f(\theta^{(s)})$ and $\sigma_f^2 = \text{Var}(X_s) < \infty$. Write $\bar{X} = S^{-1} \sum_{s=1}^S X_s$ and $\rho_{f,k} = \text{Corr}(X_s, X_{s+k})$.

- Monte Carlo recap
- Markov Chain Monte Carlo (MCMC)
 - Gibbs sampling
 - Metropolis-Hastings
- MCMC Diagnostics
 - Warm-up
 - Convergence
 - S_{eff} , MCSE, and autocorrelation
- Difficult geometries
- Additional material



From autocorrelation to Monte Carlo error

Consider one stationary chain with $X_s = f(\theta^{(s)})$ and $\sigma_f^2 = \text{Var}(X_s) < \infty$. Write $\bar{X} = S^{-1} \sum_{s=1}^S X_s$ and $\rho_{f,k} = \text{Corr}(X_s, X_{s+k})$.

$$\begin{aligned}\text{Var}(\bar{X}) &= \frac{1}{S^2} \sum_{s=1}^S \sum_{t=1}^S \text{Cov}(X_s, X_t) \\ &= \frac{\sigma_f^2}{S} \left[1 + 2 \sum_{k=1}^{S-1} \left(1 - \frac{k}{S} \right) \rho_{f,k} \right].\end{aligned}$$

Each pair contributes twice, and there are $S - k$ pairs at lag k .

- Monte Carlo recap
- Markov Chain Monte Carlo (MCMC)
 - Gibbs sampling
 - Metropolis-Hastings
- MCMC Diagnostics
 - Warm-up
 - Convergence
 - S_{eff} , MCSE, and autocorrelation
- Difficult geometries
- Additional material



From autocorrelation to Monte Carlo error

Consider one stationary chain with $X_s = f(\theta^{(s)})$ and $\sigma_f^2 = \text{Var}(X_s) < \infty$. Write $\bar{X} = S^{-1} \sum_{s=1}^S X_s$ and $\rho_{f,k} = \text{Corr}(X_s, X_{s+k})$.

$$\begin{aligned}\text{Var}(\bar{X}) &= \frac{1}{S^2} \sum_{s=1}^S \sum_{t=1}^S \text{Cov}(X_s, X_t) \\ &= \frac{\sigma_f^2}{S} \left[1 + 2 \sum_{k=1}^{S-1} \left(1 - \frac{k}{S} \right) \rho_{f,k} \right].\end{aligned}$$

Each pair contributes twice, and there are $S - k$ pairs at lag k .

For a long run with absolutely summable autocorrelations,

$$\text{Var}(\bar{X}) \approx \frac{\sigma_f^2 \tau_f}{S}, \quad \tau_f = 1 + 2 \sum_{k=1}^{\infty} \rho_{f,k}.$$

- Monte Carlo recap
- Markov Chain Monte Carlo (MCMC)
 - Gibbs sampling
 - Metropolis-Hastings
- MCMC Diagnostics
 - Warm-up
 - Convergence
 - S_{eff} , MCSE, and autocorrelation
- Difficult geometries
- Additional material



From autocorrelation to Monte Carlo error

Consider one stationary chain with $X_s = f(\theta^{(s)})$ and $\sigma_f^2 = \text{Var}(X_s) < \infty$. Write $\bar{X} = S^{-1} \sum_{s=1}^S X_s$ and $\rho_{f,k} = \text{Corr}(X_s, X_{s+k})$.

$$\begin{aligned}\text{Var}(\bar{X}) &= \frac{1}{S^2} \sum_{s=1}^S \sum_{t=1}^S \text{Cov}(X_s, X_t) \\ &= \frac{\sigma_f^2}{S} \left[1 + 2 \sum_{k=1}^{S-1} \left(1 - \frac{k}{S} \right) \rho_{f,k} \right].\end{aligned}$$

Each pair contributes twice, and there are $S - k$ pairs at lag k .

For a long run with absolutely summable autocorrelations,

$$\text{Var}(\bar{X}) \approx \frac{\sigma_f^2 \tau_f}{S}, \quad \tau_f = 1 + 2 \sum_{k=1}^{\infty} \rho_{f,k}.$$

Positive autocorrelations add uncertainty to the average, increasing its Monte Carlo standard error.

- Monte Carlo recap
- Markov Chain Monte Carlo (MCMC)
 - Gibbs sampling
 - Metropolis-Hastings
- MCMC Diagnostics
 - Warm-up
 - Convergence
 - S_{eff} , MCSE, and autocorrelation
- Difficult geometries
- Additional material



Effective sample size and MCSE

- Suppose we estimate $\mu_f = E[f(\theta) \mid y]$ using

$$\hat{\mu}_f = \frac{1}{S} \sum_{s=1}^S f(\theta^{(s)}).$$

- Monte Carlo recap
- Markov Chain Monte Carlo (MCMC)
 - Gibbs sampling
 - Metropolis-Hastings
- MCMC Diagnostics
 - Warm-up
 - Convergence
 - S_{eff} , MCSE, and autocorrelation
- Difficult geometries
- Additional material



Effective sample size and MCSE

- Suppose we estimate $\mu_f = E[f(\theta) \mid y]$ using

$$\hat{\mu}_f = \frac{1}{S} \sum_{s=1}^S f(\theta^{(s)}).$$

- If the autocorrelation time for $f(\theta)$ is τ_f , define

$$S_{\text{eff},f} = \frac{S}{\tau_f}.$$

- Monte Carlo recap
- Markov Chain Monte Carlo (MCMC)
 - Gibbs sampling
 - Metropolis-Hastings
- MCMC Diagnostics
 - Warm-up
 - Convergence
 - S_{eff} , MCSE, and autocorrelation
- Difficult geometries
- Additional material



Effective sample size and MCSE

- Suppose we estimate $\mu_f = E[f(\theta) \mid y]$ using

$$\hat{\mu}_f = \frac{1}{S} \sum_{s=1}^S f(\theta^{(s)}).$$

- If the autocorrelation time for $f(\theta)$ is τ_f , define

$$S_{\text{eff},f} = \frac{S}{\tau_f}.$$

- The Monte Carlo standard error is approximately

$$\text{MCSE}(\hat{\mu}_f) \approx \frac{\text{sd}(f(\theta))}{\sqrt{S_{\text{eff},f}}}.$$

- Monte Carlo recap
- Markov Chain Monte Carlo (MCMC)
 - Gibbs sampling
 - Metropolis-Hastings
- MCMC Diagnostics
 - Warm-up
 - Convergence
 - S_{eff} , MCSE, and autocorrelation
- Difficult geometries
- Additional material



Effective sample size and MCSE

- Suppose we estimate $\mu_f = E[f(\theta) \mid y]$ using

$$\hat{\mu}_f = \frac{1}{S} \sum_{s=1}^S f(\theta^{(s)}).$$

- If the autocorrelation time for $f(\theta)$ is τ_f , define

$$S_{\text{eff},f} = \frac{S}{\tau_f}.$$

- The Monte Carlo standard error is approximately

$$\text{MCSE}(\hat{\mu}_f) \approx \frac{\text{sd}(f(\theta))}{\sqrt{S_{\text{eff},f}}}.$$

- Larger effective sample size means a more precise Monte Carlo estimate; doubling ESS reduces MCSE by a factor of $\sqrt{2}$

- Monte Carlo recap
- Markov Chain Monte Carlo (MCMC)
 - Gibbs sampling
 - Metropolis-Hastings
- MCMC Diagnostics
 - Warm-up
 - Convergence
 - S_{eff} , MCSE, and autocorrelation
- Difficult geometries
- Additional material



Effective sample size and MCSE

- Suppose we estimate $\mu_f = E[f(\theta) \mid y]$ using

$$\hat{\mu}_f = \frac{1}{S} \sum_{s=1}^S f(\theta^{(s)}).$$

- If the autocorrelation time for $f(\theta)$ is τ_f , define

$$S_{\text{eff},f} = \frac{S}{\tau_f}.$$

- The Monte Carlo standard error is approximately

$$\text{MCSE}(\hat{\mu}_f) \approx \frac{\text{sd}(f(\theta))}{\sqrt{S_{\text{eff},f}}}.$$

- Larger effective sample size means a more precise Monte Carlo estimate; doubling ESS reduces MCSE by a factor of $\sqrt{2}$
- In practice, MCSE should be small relative to the posterior uncertainty and the precision needed for reporting

- Monte Carlo recap
- Markov Chain Monte Carlo (MCMC)
 - Gibbs sampling
 - Metropolis-Hastings
- MCMC Diagnostics
 - Warm-up
 - Convergence
 - S_{eff} , MCSE, and autocorrelation
- Difficult geometries
- Additional material



- For lag k , estimate the expected squared change within chains

$$\hat{V}_k = \frac{1}{M(N-k)} \sum_{m=1}^M \sum_{n=1}^{N-k} (\theta_{n+k,m} - \theta_{n,m})^2.$$

- Monte Carlo recap
- Markov Chain Monte Carlo (MCMC)
 - Gibbs sampling
 - Metropolis-Hastings
- MCMC Diagnostics
 - Warm-up
 - Convergence
 - S_{eff} , MCSE, and autocorrelation
- Difficult geometries
- Additional material



- For lag k , estimate the expected squared change within chains

$$\hat{V}_k = \frac{1}{M(N-k)} \sum_{m=1}^M \sum_{n=1}^{N-k} (\theta_{n+k,m} - \theta_{n,m})^2.$$

- Estimate the lag- k autocorrelation by

$$\hat{\rho}_k = 1 - \frac{\hat{V}_k}{2\widehat{\text{var}}^+}.$$

- Monte Carlo recap
- Markov Chain Monte Carlo (MCMC)
 - Gibbs sampling
 - Metropolis-Hastings
- MCMC Diagnostics
 - Warm-up
 - Convergence
 - S_{eff} , MCSE, and autocorrelation
- Difficult geometries
- Additional material



- For lag k , estimate the expected squared change within chains

$$\hat{V}_k = \frac{1}{M(N-k)} \sum_{m=1}^M \sum_{n=1}^{N-k} (\theta_{n+k,m} - \theta_{n,m})^2.$$

- Estimate the lag- k autocorrelation by

$$\hat{\rho}_k = 1 - \frac{\hat{V}_k}{2\widehat{\text{var}}^+}.$$

- Small squared changes imply high autocorrelation; large changes imply lower autocorrelation

- Monte Carlo recap
- Markov Chain Monte Carlo (MCMC)
 - Gibbs sampling
 - Metropolis-Hastings
- MCMC Diagnostics
 - Warm-up
 - Convergence
 - S_{eff} , MCSE, and autocorrelation
- Difficult geometries
- Additional material



- For lag k , estimate the expected squared change within chains

$$\hat{V}_k = \frac{1}{M(N-k)} \sum_{m=1}^M \sum_{n=1}^{N-k} (\theta_{n+k,m} - \theta_{n,m})^2.$$

- Estimate the lag- k autocorrelation by

$$\hat{\rho}_k = 1 - \frac{\hat{V}_k}{2\widehat{\text{var}}^+}.$$

- Small squared changes imply high autocorrelation; large changes imply lower autocorrelation
- Using $\widehat{\text{var}}^+$ incorporates information from all chains and penalizes chains that have not mixed well

- Monte Carlo recap
- Markov Chain Monte Carlo (MCMC)
 - Gibbs sampling
 - Metropolis-Hastings
- MCMC Diagnostics
 - Warm-up
 - Convergence
 - S_{eff} , MCSE, and autocorrelation
- Difficult geometries
- Additional material



Integrated autocorrelation time

- Let ρ_k be the autocorrelation at lag k . The integrated autocorrelation time is

$$\tau = 1 + 2 \sum_{k=1}^{\infty} \rho_k, \quad S_{\text{eff}} = \frac{S}{\tau}.$$

- Monte Carlo recap
- Markov Chain Monte Carlo (MCMC)
 - Gibbs sampling
 - Metropolis-Hastings
- MCMC Diagnostics
 - Warm-up
 - Convergence
 - S_{eff} , MCSE, and autocorrelation
- Difficult geometries
- Additional material



Integrated autocorrelation time

- Let ρ_k be the autocorrelation at lag k . The integrated autocorrelation time is

$$\tau = 1 + 2 \sum_{k=1}^{\infty} \rho_k, \quad S_{\text{eff}} = \frac{S}{\tau}.$$

- τ measures how many correlated draws provide roughly the same information as one independent draw

- Monte Carlo recap
- Markov Chain Monte Carlo (MCMC)
 - Gibbs sampling
 - Metropolis-Hastings
- MCMC Diagnostics
 - Warm-up
 - Convergence
 - S_{eff} , MCSE, and autocorrelation
- Difficult geometries
- Additional material



Integrated autocorrelation time

- Let ρ_k be the autocorrelation at lag k . The integrated autocorrelation time is

$$\tau = 1 + 2 \sum_{k=1}^{\infty} \rho_k, \quad S_{\text{eff}} = \frac{S}{\tau}.$$

- τ measures how many correlated draws provide roughly the same information as one independent draw
- Estimated autocorrelations become noisy at long lags
 - practical estimators truncate the sum at a finite lag K

$$\hat{\tau} = 1 + 2 \sum_{k=1}^K \hat{\rho}_k$$

- statistical software chooses the maximum lag K adaptively

- Monte Carlo recap
- Markov Chain Monte Carlo (MCMC)
 - Gibbs sampling
 - Metropolis-Hastings
- MCMC Diagnostics
 - Warm-up
 - Convergence
 - S_{eff} , MCSE, and autocorrelation
- Difficult geometries
- Additional material



Fast ACF decay and larger ESS




- All examples use $S = 4500$ retained draws after warm-up
- Pair discussion: rank the three chains by expected ESS using their ACF plots
- Which chain should give the smallest MCSE if their posterior standard deviations are similar?

- Monte Carlo recap
- Markov Chain Monte Carlo (MCMC)
 - Gibbs sampling
 - Metropolis-Hastings
- MCMC Diagnostics
 - Warm-up
 - Convergence
 - S_{eff} , MCSE, and autocorrelation
- Difficult geometries
- Additional material



Fast ACF decay and larger ESS

Fastest decay: $\hat{\tau} \approx 24 \Rightarrow \hat{S}_{\text{eff}} \approx 188$ for $S = 4500$ retained draws



— theta1 — theta2 - - 95% interval for MCMC error ··· 95% interval for independent

ESS for one scalar parameter:

$$\hat{\tau} = 1 + 2 \sum_{k=1}^K \hat{\rho}_k$$

$$\approx 24$$

$$\hat{S}_{\text{eff}} = \frac{4500}{24} \approx 188$$




- Monte Carlo recap
- Markov Chain Monte Carlo (MCMC)
 - Gibbs sampling
 - Metropolis-Hastings
- MCMC Diagnostics
 - Warm-up
 - Convergence
 - S_{eff} , MCSE, and autocorrelation
- Difficult geometries
- Additional material



- Monte Carlo recap
- Markov Chain Monte Carlo (MCMC)
 - Gibbs sampling
 - Metropolis-Hastings
- MCMC Diagnostics
 - Warm-up
 - Convergence
 - S_{eff} , MCSE, and autocorrelation
- Difficult geometries
- Additional material

Slow ACF decay and smaller ESS

Slowest ACF decay: $\hat{\tau} \approx 104$ and $\hat{S}_{\text{eff}} \approx 43$ from $S = 4500$ retained draws



— theta1 — theta2 - - 95% interval for MCMC error ... 95% interval for independent

ESS for one scalar parameter:

$$\hat{\tau} = 1 + 2 \sum_{k=1}^K \hat{\rho}_k$$

$$\approx 104$$




$$\hat{S}_{\text{eff}} = \frac{4500}{104} \approx 43$$



- Monte Carlo recap
- Markov Chain Monte Carlo (MCMC)
 - Gibbs sampling
 - Metropolis-Hastings
- MCMC Diagnostics
 - Warm-up
 - Convergence
 - S_{eff} , MCSE, and autocorrelation
- Difficult geometries
- Additional material

Comparing autocorrelation and ESS

Across all three examples: slower ACF decay $\Rightarrow$ larger $\hat{\tau} \Rightarrow$ smaller $\hat{S}_{\text{eff}}$.



— theta1 — theta2 — 95% interval for MCMC error — 95% interval for independent

ESS for one scalar parameter:

$$\hat{\tau} = 1 + 2 \sum_{k=1}^K \hat{\rho}_k$$

$$\approx 63$$

$$\hat{S}_{\text{eff}} = \frac{4500}{63} \approx 71$$



UPPSALA
UNIVERSITET

- Monte Carlo recap
- Markov Chain Monte Carlo (MCMC)
 - Gibbs sampling
 - Metropolis-Hastings
- MCMC Diagnostics
 - Warm-up
 - Convergence
 - S_{eff} , MCSE, and autocorrelation
- Difficult geometries
- Additional material

Section 4

Difficult geometries



Difficult geometry: changing scale and direction

- Monte Carlo recap
- Markov Chain Monte Carlo (MCMC)
 - Gibbs sampling
 - Metropolis-Hastings
- MCMC Diagnostics
 - Warm-up
 - Convergence
 - S_{eff} , MCSE, and autocorrelation
- Difficult geometries
- Additional material

- **Nonlinear dependence:** posterior mass follows a curved, narrow region
 - a fixed proposal shape and scale cannot follow the curve, so many proposals fall outside the high-density region
 - symptoms: slow movement, high autocorrelation, and low ESS



- Monte Carlo recap
- Markov Chain Monte Carlo (MCMC)
 - Gibbs sampling
 - Metropolis-Hastings
- MCMC Diagnostics
 - Warm-up
 - Convergence
 - S_{eff} , MCSE, and autocorrelation
- Difficult geometries
- Additional material

Difficult geometry: changing scale and direction

- **Nonlinear dependence:** posterior mass follows a curved, narrow region
 - a fixed proposal shape and scale cannot follow the curve, so many proposals fall outside the high-density region
 - symptoms: slow movement, high autocorrelation, and low ESS
- **Funnel geometry:** the appropriate step size changes across the posterior
 - a step size suitable for the wide part is too large for the neck
 - a step size suitable for the neck moves very slowly in the wide part
 - symptoms: chains explore different scales or rarely enter the neck

Banana demo

Funnel demo



- **Multimodality:** chains can remain in different modes
 - chain means disagree and $\hat{R}$ can remain above 1
 - if every chain starts in the same mode, diagnostics may fail to reveal the missing mode

- Monte Carlo recap
- Markov Chain Monte Carlo (MCMC)
 - Gibbs sampling
 - Metropolis-Hastings
- MCMC Diagnostics
 - Warm-up
 - Convergence
 - S_{eff} , MCSE, and autocorrelation
- Difficult geometries
- Additional material



- **Multimodality:** chains can remain in different modes
 - chain means disagree and $\hat{R}$ can remain above 1
 - if every chain starts in the same mode, diagnostics may fail to reveal the missing mode
- **Non-identifiability:** many parameter combinations fit the data almost equally well
 - with weak prior information, the posterior can contain a long ridge
 - symptoms: drifting traces, strong dependence, and low ESS

- Monte Carlo recap
- Markov Chain Monte Carlo (MCMC)
 - Gibbs sampling
 - Metropolis-Hastings
- MCMC Diagnostics
 - Warm-up
 - Convergence
 - S_{eff} , MCSE, and autocorrelation
- Difficult geometries
- Additional material



- Monte Carlo recap
- Markov Chain Monte Carlo (MCMC)
 - Gibbs sampling
 - Metropolis-Hastings
- MCMC Diagnostics
 - Warm-up
 - Convergence
 - S_{eff} , MCSE, and autocorrelation
- Difficult geometries
- Additional material

Difficult geometry: hidden regions

- **Multimodality**: chains can remain in different modes
 - chain means disagree and $\hat{R}$ can remain above 1
 - if every chain starts in the same mode, diagnostics may fail to reveal the missing mode
- **Non-identifiability**: many parameter combinations fit the data almost equally well
 - with weak prior information, the posterior can contain a long ridge
 - symptoms: drifting traces, strong dependence, and low ESS
- **Heavy tails**: occasional extreme draws dominate summaries
 - means or variances may not exist, so the usual CLT and MCSE may not apply

Multimodal demo

Ridge demo



UPPSALA UNIVERSITET

- Monte Carlo recap
- Markov Chain Monte Carlo (MCMC)
 - Gibbs sampling
 - Metropolis-Hastings
- MCMC Diagnostics
 - Warm-up
 - Convergence
 - S_{eff} , MCSE, and autocorrelation
- Difficult geometries
- Additional material

MCMC diagnostic workflow

1. Run several chains from different, reasonably dispersed starting values
2. Remove warm-up draws before computing posterior summaries
3. Inspect trace plots: chains should explore the same regions without persistent trends
4. Check that $\hat{R} < 1.01$ for all reported quantities
5. Check ESS and MCSE for the posterior summaries that matter
6. If diagnostics are poor, investigate the proposal, parameterization, model, and posterior geometry



UPPSALA UNIVERSITET

- Monte Carlo recap
- Markov Chain Monte Carlo (MCMC)
 - Gibbs sampling
 - Metropolis-Hastings
- MCMC Diagnostics
 - Warm-up
 - Convergence
 - S_{eff} , MCSE, and autocorrelation
- Difficult geometries
- Additional material

MCMC diagnostic workflow

1. Run several chains from different, reasonably dispersed starting values
2. Remove warm-up draws before computing posterior summaries
3. Inspect trace plots: chains should explore the same regions without persistent trends
4. Check that $\hat{R} < 1.01$ for all reported quantities
5. Check ESS and MCSE for the posterior summaries that matter
6. If diagnostics are poor, investigate the proposal, parameterization, model, and posterior geometry

Running longer helps only if the chain is exploring the relevant posterior regions.



UPPSALA
UNIVERSITET

- Monte Carlo recap
- Markov Chain Monte Carlo (MCMC)
 - Gibbs sampling
 - Metropolis-Hastings
- MCMC Diagnostics
 - Warm-up
 - Convergence
 - S_{eff} , MCSE, and autocorrelation
- Difficult geometries
- Additional material

Section 5

Additional material



- Monte Carlo recap
 - Markov Chain Monte Carlo (MCMC)
 - Gibbs sampling
 - Metropolis-Hastings
 - MCMC Diagnostics
 - Warm-up
 - Convergence
 - S_{eff} , MCSE, and autocorrelation
 - Difficult geometries
 - **Additional material**
- BUGS popularized automatic Gibbs sampling for Bayesian models
 - WinBUGS and OpenBUGS provided widely used BUGS implementations
 - JAGS follows a similar model-description and sampling approach
 - These systems work especially naturally for models with conditionally conjugate components



UPPSALA
UNIVERSITET

$\hat{R}$

- M chains, each having N draws

- Monte Carlo recap
- Markov Chain Monte Carlo (MCMC)
 - Gibbs sampling
 - Metropolis-Hastings
- MCMC Diagnostics
 - Warm-up
 - Convergence
 - S_{eff} , MCSE, and autocorrelation
- Difficult geometries
- Additional material



- Monte Carlo recap
- Markov Chain Monte Carlo (MCMC)
 - Gibbs sampling
 - Metropolis-Hastings
- MCMC Diagnostics
 - Warm-up
 - Convergence
 - S_{eff} , MCSE, and autocorrelation
- Difficult geometries
- Additional material

- M chains, each having N draws
- Within chains variance W

$$W = \frac{1}{M} \sum_{m=1}^M s_m^2, \text{ where } s_m^2 = \frac{1}{N-1} \sum_{n=1}^N (\theta_{nm} - \bar{\theta}_{.m})^2$$



- Monte Carlo recap
- Markov Chain Monte Carlo (MCMC)
 - Gibbs sampling
 - Metropolis-Hastings
- MCMC Diagnostics
 - Warm-up
 - Convergence
 - S_{eff} , MCSE, and autocorrelation
- Difficult geometries
- Additional material

- M chains, each having N draws
- Within chains variance W

$$W = \frac{1}{M} \sum_{m=1}^M s_m^2, \text{ where } s_m^2 = \frac{1}{N-1} \sum_{n=1}^N (\theta_{nm} - \bar{\theta}_{.m})^2$$

- Between chains variance B

$$B = \frac{N}{M-1} \sum_{m=1}^M (\bar{\theta}_{.m} - \bar{\theta}_{..})^2,$$

$$\text{where } \bar{\theta}_{.m} = \frac{1}{N} \sum_{n=1}^N \theta_{nm}, \bar{\theta}_{..} = \frac{1}{M} \sum_{m=1}^M \bar{\theta}_{.m}$$



- Monte Carlo recap
- Markov Chain Monte Carlo (MCMC)
 - Gibbs sampling
 - Metropolis-Hastings
- MCMC Diagnostics
 - Warm-up
 - Convergence
 - S_{eff} , MCSE, and autocorrelation
- Difficult geometries
- Additional material

- M chains, each having N draws
- Within chains variance W

$$W = \frac{1}{M} \sum_{m=1}^M s_m^2, \text{ where } s_m^2 = \frac{1}{N-1} \sum_{n=1}^N (\theta_{nm} - \bar{\theta}_{.m})^2$$

- Between chains variance B

$$B = \frac{N}{M-1} \sum_{m=1}^M (\bar{\theta}_{.m} - \bar{\theta}_{..})^2,$$

$$\text{where } \bar{\theta}_{.m} = \frac{1}{N} \sum_{n=1}^N \theta_{nm}, \bar{\theta}_{..} = \frac{1}{M} \sum_{m=1}^M \bar{\theta}_{.m}$$

- B/N is variance of the means of the chains



- Monte Carlo recap
- Markov Chain Monte Carlo (MCMC)
 - Gibbs sampling
 - Metropolis-Hastings
- MCMC Diagnostics
 - Warm-up
 - Convergence
 - S_{eff} , MCSE, and autocorrelation
- Difficult geometries
- Additional material

- M chains, each having N draws
- Within chains variance W

$$W = \frac{1}{M} \sum_{m=1}^M s_m^2, \text{ where } s_m^2 = \frac{1}{N-1} \sum_{n=1}^N (\theta_{nm} - \bar{\theta}_{.m})^2$$

- Between chains variance B

$$B = \frac{N}{M-1} \sum_{m=1}^M (\bar{\theta}_{.m} - \bar{\theta}_{..})^2,$$

$$\text{where } \bar{\theta}_{.m} = \frac{1}{N} \sum_{n=1}^N \theta_{nm}, \bar{\theta}_{..} = \frac{1}{M} \sum_{m=1}^M \bar{\theta}_{.m}$$

- B/N is variance of the means of the chains
- Estimate total variance $\text{var}(\theta|y)$ as a weighted mean of W and B

$$\widehat{\text{var}}^+(\theta|y) = \frac{N-1}{N} W + \frac{1}{N} B$$



- Monte Carlo recap
- Markov Chain Monte Carlo (MCMC)
 - Gibbs sampling
 - Metropolis-Hastings
- MCMC Diagnostics
 - Warm-up
 - Convergence
 - S_{eff} , MCSE, and autocorrelation
- Difficult geometries
- Additional material

- Estimate total variance $\text{var}(\theta|y)$ as a weighted mean of W and B

$$\widehat{\text{var}}^+(\theta|y) = \frac{N-1}{N}W + \frac{1}{N}B$$

- with overdispersed starting points and incomplete mixing, this typically *overestimates* the marginal posterior variance



- Monte Carlo recap
- Markov Chain Monte Carlo (MCMC)
 - Gibbs sampling
 - Metropolis-Hastings
- MCMC Diagnostics
 - Warm-up
 - Convergence
 - S_{eff} , MCSE, and autocorrelation
- Difficult geometries
- **Additional material**

- Estimate total variance $\text{var}(\theta|y)$ as a weighted mean of W and B

$$\widehat{\text{var}}^+(\theta|y) = \frac{N-1}{N}W + \frac{1}{N}B$$

- with overdispersed starting points and incomplete mixing, this typically *overestimates* the marginal posterior variance
- Before chains have fully explored the posterior, W typically *underestimates* the marginal posterior variance
 - single chains have not yet visited all points in the distribution
 - when $N \rightarrow \infty$, $E(W) \rightarrow \text{var}(\theta|y)$



- Monte Carlo recap
- Markov Chain Monte Carlo (MCMC)
 - Gibbs sampling
 - Metropolis-Hastings
- MCMC Diagnostics
 - Warm-up
 - Convergence
 - S_{eff} , MCSE, and autocorrelation
- Difficult geometries
- **Additional material**

- Estimate total variance $\text{var}(\theta|y)$ as a weighted mean of W and B

$$\widehat{\text{var}}^+(\theta|y) = \frac{N-1}{N}W + \frac{1}{N}B$$

- with overdispersed starting points and incomplete mixing, this typically *overestimates* the marginal posterior variance
- Before chains have fully explored the posterior, W typically *underestimates* the marginal posterior variance
 - single chains have not yet visited all points in the distribution
 - when $N \rightarrow \infty$, $E(W) \rightarrow \text{var}(\theta|y)$
- Comparing these two estimates motivates

$$\hat{R} = \sqrt{\frac{\widehat{\text{var}}^+}{W}}$$



Geyer's adaptive window estimator of τ

- The maximum lag K can be chosen adaptively
 - for a stationary, reversible Markov chain
 - let $\Gamma_m = \rho_{2m} + \rho_{2m+1}$ for $m = 0, 1, \dots$ and $\rho_0 = 1$
 - Γ_m is a positive, decreasing, and convex function of m

- Monte Carlo recap
- Markov Chain Monte Carlo (MCMC)
 - Gibbs sampling
 - Metropolis-Hastings
- MCMC Diagnostics
 - Warm-up
 - Convergence
 - S_{eff} , MCSE, and autocorrelation
- Difficult geometries
- Additional material



- Monte Carlo recap
- Markov Chain Monte Carlo (MCMC)
 - Gibbs sampling
 - Metropolis-Hastings
- MCMC Diagnostics
 - Warm-up
 - Convergence
 - S_{eff} , MCSE, and autocorrelation
- Difficult geometries
- Additional material

Geyer's adaptive window estimator of τ

- The maximum lag K can be chosen adaptively
 - for a stationary, reversible Markov chain
 - let $\Gamma_m = \rho_{2m} + \rho_{2m+1}$ for $m = 0, 1, \dots$ and $\rho_0 = 1$
 - Γ_m is a positive, decreasing, and convex function of m
- Initial positive sequence estimator (Geyer's IPSE)
 - Starting at $m = 0$, retain consecutive positive estimated pairs and stop before the first nonpositive pair
 - If L is the last retained pair, the maximum lag is $K = 2L + 1$

$$\hat{\tau} = -1 + 2 \sum_{m=0}^L \hat{\Gamma}_m = 1 + 2 \sum_{k=1}^K \hat{\rho}_k.$$

