



UPPSALA
UNIVERSITET

- Hierarchical models
 - Rats example
 - Factory example
 - 8 schools example
- Exchangeability
- Computational aspects

Bayesian Statistics and Data Analysis

Lecture 7

Måns Magnusson

Department of Statistics, Uppsala University
Thanks to Aki Vehtari, Aalto University



UPPSALA
UNIVERSITET

- **Hierarchical models**
 - Rats example
 - Factory example
 - 8 schools example
- Exchangeability
- Computational aspects

Section 1

Hierarchical models



UPPSALA
UNIVERSITET

Stan tutorial for hierarchical models

- Hierarchical models

- Rats example
- Factory example
- 8 schools example

- Exchangeability

- Computational aspects

- Before Assignment 7, watch [Hierarchical Modelling in Stan: Predicting the Premier League](#)
- While watching, connect the example to:
 - partial pooling
 - group-level parameters
 - comparison with non-hierarchical models
 - Stan diagnostics for hierarchical models



- **Hierarchical models**


- Rats example
- Factory example
- 8 schools example

- Exchangeability

- Computational aspects

Survival probabilities across hospitals

- Example: comparing hospitals
 - Hospital j : survival probability θ_j
 - Patient i 's survival outcome: y_{ij}





- **Hierarchical models**


- Rats example
- Factory example
- 8 schools example

- **Exchangeability**


- **Computational aspects**

Survival probabilities across hospitals

- Example: comparing hospitals
 - Hospital j : survival probability θ_j
 - Patient i 's survival outcome: y_{ij}



- Model the θ_j as similar, not identical



- A shared population distribution has hyperparameters ϕ
- Both θ_j and ϕ are unknown



- **Hierarchical models**

- Rats example
- Factory example
- 8 schools example

- Exchangeability

- Computational aspects


- $j = 1, \dots, J$: groups; $i = 1, \dots, n_j$: observations in group j
- y_{ij} : individual observation; y_j : group data or a group summary; $y = (y_1, \dots, y_J)$
- θ_j : group-specific parameter; $\theta = (\theta_1, \dots, \theta_J)$
- ϕ : population hyperparameters

For normal hierarchies:

- μ : population mean; τ : between-group SD
- σ : within-group SD of individual observations




- **Hierarchical models**
 - Rats example
 - Factory example
 - 8 schools example
- Exchangeability
- Computational aspects



- Given ϕ , the θ_j are independent with a common population distribution $p(\theta_j | \phi)$



- **Hierarchical models**
 - Rats example
 - Factory example
 - 8 schools example
- Exchangeability
- Computational aspects



- Given ϕ , the θ_j are independent with a common population distribution $p(\theta_j | \phi)$
- Given θ , groups are independent, and y_j depends only on its own parameter θ_j
- The hyperprior $p(\phi)$ describes uncertainty about the population distribution



Hierarchical model: joint posterior

Bayes' rule and the conditional-independence assumptions give

$$\begin{aligned} p(\theta, \phi \mid y) &\propto p(y \mid \theta) p(\theta \mid \phi) p(\phi) \\ &\propto p(\phi) \prod_{j=1}^J [p(\theta_j \mid \phi) p(y_j \mid \theta_j)] . \end{aligned}$$

- One hyperprior; each group contributes a population factor and a likelihood

- Hierarchical models
 - Rats example
 - Factory example
 - 8 schools example
- Exchangeability
- Computational aspects



Hierarchical model: joint posterior

Bayes' rule and the conditional-independence assumptions give

$$\begin{aligned} p(\theta, \phi \mid y) &\propto p(y \mid \theta) p(\theta \mid \phi) p(\phi) \\ &\propto p(\phi) \prod_{j=1}^J [p(\theta_j \mid \phi) p(y_j \mid \theta_j)] . \end{aligned}$$

- Hierarchical models
 - Rats example
 - Factory example
 - 8 schools example
- Exchangeability
- Computational aspects
- One hyperprior; each group contributes a population factor and a likelihood
- Given ϕ and y , the group parameters are independent
- Averaging over the shared, uncertain ϕ generally makes the θ_j dependent in the posterior
- Each group's data inform ϕ , so other groups can inform its estimate of θ_j




- **Hierarchical models**
 - Rats example
 - Factory example
 - 8 schools example
- Exchangeability
- Computational aspects

No, complete, and partial pooling

No pooling (separate)

Independent priors; each θ_j uses only group j 's data.





- **Hierarchical models**

- Rats example
- Factory example
- 8 schools example


- Exchangeability

- Computational aspects

No, complete, and partial pooling


No pooling (separate)

Independent priors; each θ_j uses only group j 's data.



Complete pooling (pooled)

All groups inform one common θ .





- **Hierarchical models**


- Rats example
- Factory example
- 8 schools example

- Exchangeability
- Computational aspects

No, complete, and partial pooling


No pooling (separate)

Independent priors; each θ_j uses only group j 's data.




Complete pooling (pooled)

All groups inform one common θ .



Partial pooling (hierarchical)

Distinct θ_j share information through ϕ .





- Hierarchical models
 - Rats example
 - Factory example
 - 8 schools example
- Exchangeability
- Computational aspects

Pair discussion: group size and pooling

Discuss in pairs (1–2 minutes).

Two groups have the same observed success rate:

$$\text{Group A: } \frac{1}{5} = 20\%, \quad \text{Group B: } \frac{100}{500} = 20\%.$$

Both share the same population information, with a population mean different from 20%.

Which estimate should move more toward the population mean? Why?



Pair discussion: group size and pooling

Discuss in pairs (1–2 minutes).

Two groups have the same observed success rate:

$$\text{Group A: } \frac{1}{5} = 20\%, \quad \text{Group B: } \frac{100}{500} = 20\%.$$

Both share the same population information, with a population mean different from 20%.

Which estimate should move more toward the population mean? Why?

Answer: Group A. Its five observations provide weaker evidence, so the population information receives more weight.

- Hierarchical models
 - Rats example
 - Factory example
 - 8 schools example
- Exchangeability
- Computational aspects



Predictive distributions for hierarchical models


- Hierarchical models

- Rats example
- Factory example
- 8 schools example

- Exchangeability

- Computational aspects

- Two types of predictive distributions
 1. A new observation in **an existing group**
 2. A new observation in **a new group**





UPPSALA
UNIVERSITET

Hierarchical binomial model: rats

- Hierarchical models
 - Rats example
 - Factory example
 - 8 schools example
 - Exchangeability
 - Computational aspects
- Medicine testing
 - Type F344 female rats in control group given placebo
 - count how many get endometrial stromal polyps
 - familiar binomial model example



- Hierarchical models
 - Rats example
 - Factory example
 - 8 schools example
- Exchangeability
- Computational aspects

- Medicine testing
- Type F344 female rats in control group given placebo
 - count how many get endometrial stromal polyps
 - familiar binomial model example
- Experiment has been repeated 71 times


0/20	0/20	0/20	0/20	0/20	0/20	0/20	0/19	0/19	0/19
0/19	0/18	0/18	0/17	1/20	1/20	1/20	1/20	1/19	1/19
1/18	1/18	2/25	2/24	2/23	2/20	2/20	2/20	2/20	2/20
2/20	1/10	5/49	2/19	5/46	3/27	2/17	7/49	7/47	3/20
3/20	2/13	9/48	10/50	4/20	4/20	4/20	4/20	4/20	4/20
4/20	10/48	4/19	4/19	4/19	5/22	11/46	12/49	5/20	5/20
6/23	5/19	6/22	6/20	6/20	6/20	16/52	15/46	15/47	9/24
4/14									



Hierarchical binomial model: rats

What assumption does each model make about the experiments?

Pooled model



Complete pooling assumes one common probability for all experiments.




Hierarchical binomial model: rats

What assumption does each model make about the experiments?

Separate model

Existing-group posteriors




No pooling estimates each probability separately, without learning a shared population distribution.



- Hierarchical models
 - Rats example
 - Factory example
 - 8 schools example
- Exchangeability
- Computational aspects

Reading the model diagrams



- Circle: unknown quantity
- Single box: observed quantity or a supplied estimate
- Double box: fixed, given quantity
- Arrow: the child's distribution depends on the parent
- Outer box (**plate**): repeat its contents for $j = 1, \dots, J$


These are the conventions used in this lecture. A nested i plate repeats observations within each group.



- Hierarchical models
 - Rats example
 - Factory example
 - 8 schools example
- Exchangeability
- Computational aspects

Hierarchical binomial model: rats

- Here $\phi = (\alpha, \beta)$; y_j counts affected rats out of known n_j



$$\theta_j \mid \alpha, \beta \sim \text{Beta}(\alpha, \beta)$$


$$y_j \mid \theta_j, n_j \sim \text{Binomial}(n_j, \theta_j)$$



- Hierarchical models
 - Rats example
 - Factory example
 - 8 schools example
- Exchangeability
- Computational aspects

Hierarchical binomial model: rats

- Here $\phi = (\alpha, \beta)$; y_j counts affected rats out of known n_j



$$\theta_j \mid \alpha, \beta \sim \text{Beta}(\alpha, \beta)$$

$$y_j \mid \theta_j, n_j \sim \text{Binomial}(n_j, \theta_j)$$

$$p(\theta, \alpha, \beta \mid y) = \left[\prod_{j=1}^J p(\theta_j \mid \alpha, \beta, y_j) \right] p(\alpha, \beta \mid y)$$

- One hyperposterior, not one per group
- The θ_j are independent given α, β, y , but generally dependent after averaging over α, β



For fixed $\alpha, \beta > 0$, the binomial likelihood gives

$$\begin{aligned} p(\theta_j \mid \alpha, \beta, y_j) &\propto \theta_j^{y_j} (1 - \theta_j)^{n_j - y_j} \theta_j^{\alpha - 1} (1 - \theta_j)^{\beta - 1} \\ &\propto \theta_j^{\alpha + y_j - 1} (1 - \theta_j)^{\beta + n_j - y_j - 1}. \end{aligned}$$

Thus

$$\theta_j \mid \alpha, \beta, y_j \sim \text{Beta}(\alpha + y_j, \beta + n_j - y_j).$$

- The population distribution supplies α, β
- Group j supplies y_j successes and $n_j - y_j$ failures

- Hierarchical models
 - Rats example
 - Factory example
 - 8 schools example
- Exchangeability
- Computational aspects



Beta-binomial shrinkage

Write the population mean and concentration as

$$m = \frac{\alpha}{\alpha + \beta}, \quad \kappa = \alpha + \beta.$$

For $n_j > 0$,

$$\begin{aligned} E(\theta_j \mid \alpha, \beta, y_j) &= \frac{\alpha + y_j}{\alpha + \beta + n_j} \\ &= \underbrace{\frac{n_j}{n_j + \kappa}}_{\text{data weight}} \frac{y_j}{n_j} + \underbrace{\frac{\kappa}{n_j + \kappa}}_{\text{population weight}} m. \end{aligned}$$

- The conditional posterior mean lies between y_j/n_j and m
- With α, β fixed, smaller groups receive more population weight
- Full Bayes also averages over uncertainty in α, β



- Hierarchical models
 - Rats example
 - Factory example
 - 8 schools example
- Exchangeability
- Computational aspects

- Population prior $\text{Beta}(\theta_j | \alpha, \beta)$
- Hyperprior $p(\alpha, \beta)$?
 - α, β both affect the location and scale
 - BDA3 (pp. 110–111) uses, for $\alpha, \beta > 0$,
 $p(\alpha, \beta) \propto (\alpha + \beta)^{-5/2}$
 - This hyperprior is **improper**: its integral is infinite
 - BDA shows that the posterior is proper for these rat data; posterior propriety must be checked when using an improper prior
 - An example-specific choice, not a general default



- Hierarchical models
 - Rats example
 - Factory example
 - 8 schools example
- Exchangeability
- Computational aspects

- Population prior $\text{Beta}(\theta_j | \alpha, \beta)$
- Hyperprior $p(\alpha, \beta)$?
 - α, β both affect the location and scale
 - BDA3 (pp. 110–111) uses, for $\alpha, \beta > 0$,
 $p(\alpha, \beta) \propto (\alpha + \beta)^{-5/2}$
 - This hyperprior is **improper**: its integral is infinite
 - BDA shows that the posterior is proper for these rat data; posterior propriety must be checked when using an improper prior
 - An example-specific choice, not a general default
- What predictive distributions can we have?



- Hierarchical models
 - Rats example
 - Factory example
 - 8 schools example
- Exchangeability
- Computational aspects

Hierarchical inference with posterior draws

For $s = 1, \dots, S$:

1. Draw $(\alpha^{(s)}, \beta^{(s)}) \sim p(\alpha, \beta \mid y)$
2. For every group j , draw independently conditional on this pair:

$$\theta_j^{(s)} \sim \text{Beta}(\alpha^{(s)} + y_j, \beta^{(s)} + n_j - y_j)$$

Share one hyperparameter draw across all groups at each step.

$$\mathbb{E}[h(\theta_j) \mid y] = \int h(\theta_j) p(\theta_j \mid y) d\theta_j \approx \frac{1}{S} \sum_{s=1}^S h(\theta_j^{(s)}).$$


Choosing $h(\theta) = \theta$ estimates a mean; $h(\theta) = \mathbf{1}(\theta > c)$ estimates a probability.



Rats: learning the population distribution

Which population distributions remain plausible?

The marginal of α and β



The hyperposterior $p(\alpha, \beta \mid y)$ uses all 71 experiments.




- Hierarchical models
 - Rats example
 - Factory example
 - 8 schools example
- Exchangeability
- Computational aspects

Rats: learning the population distribution

How does hyperparameter uncertainty affect the population?

Beta(α, β) given posterior draws of α and β



Each curve is Beta(α, β) for one posterior draw of (α, β) .




Rats: a new experiment's probability

What does the green curve predict?

New group's probability

Posterior predictive $p(\theta_{\text{new}} | y)$



$$p(\theta_{\text{new}} | y) = \int_0^{\infty} \int_0^{\infty} p(\theta_{\text{new}} | \alpha, \beta) p(\alpha, \beta | y) d\alpha d\beta.$$

This averages the population curves over the hyperposterior. It predicts a new experiment's probability, before observing its data.




- Hierarchical models
 - Rats example
 - Factory example
 - 8 schools example
- Exchangeability
- Computational aspects

Rats: partial pooling

Follow experiment 71 (4/14), shown in red. What changes?


Separate model

Existing-group posteriors



Hierarchical model

Existing-group posteriors



Experiment 71 moves toward the population.
Both plots use the same experiments and scales.



- Hierarchical models

- Rats example
- Factory example
- 8 schools example

- Exchangeability

- Computational aspects


Rats: existing and new experiments

Why do these distributions answer different questions?

Hierarchical model


Existing-group posteriors

Experiment 71 (4/14)
Other experiments



New group's probability

Posterior predictive $p(\theta_{\text{new}} | y)$



Existing $p(\theta_j | y)$ uses group j 's data as well as the population.
New $p(\theta_{\text{new}} | y)$ has no new-group data.




- Hierarchical models
 - Rats example
 - **Factory example**
 - 8 schools example
- Exchangeability
- Computational aspects

Hierarchical normal model: factory

- We evaluate the quality produced by six machines in a factory
- Assume hierarchical model
 - each machine has its own (average) quality θ_j and **common variance** σ^2

$$\theta_j | \mu, \tau \sim \mathcal{N}(\mu, \tau^2)$$

$$y_{ij} | \theta_j \sim \mathcal{N}(\theta_j, \sigma^2)$$



- Can be used to predict the future quality produced by each machine and quality produced by a new similar machine



- Hierarchical models
 - Rats example
 - **Factory example**
 - 8 schools example
- Exchangeability
- Computational aspects

Factory model: a complete hierarchy

One illustrative choice of **proper, independent priors** is

$$\begin{aligned}\mu &\sim \mathcal{N}(m_0, s_0^2), \\ \tau &\sim \text{HalfNormal}(A_\tau), \quad \sigma \sim \text{HalfNormal}(A_\sigma).\end{aligned}$$

The common-variance factory model then has

$$\begin{aligned}\theta_j &| \mu, \tau \stackrel{\text{iid}}{\sim} \mathcal{N}(\mu, \tau^2), & j = 1, \dots, J, \\ y_{ij} &| \theta_j, \sigma \sim \mathcal{N}(\theta_j, \sigma^2), & i = 1, \dots, n_j.\end{aligned}$$

- Observations are independent given the group means and σ
- Fix $m_0, s_0, A_\tau, A_\sigma$ using subject knowledge and the measurement scale
- Here $\mathcal{N}$ uses variance as its second argument; A_τ, A_σ are half-normal scale parameters



A half-normal prior for a standard deviation

A simple way to restrict τ to positive values is

$$Z \sim \mathcal{N}(0, A_\tau^2), \quad \tau = |Z|.$$

This defines $\tau \sim \text{HalfNormal}(A_\tau)$, with

$$\Pr(\tau < 1.96A_\tau) \approx 0.95.$$

- A_τ has the same units as the outcome
- Example: to put 95% prior probability below a between-machine SD of 20 quality units, choose $A_\tau \approx 20/1.96 = 10.2$
- Choose A_σ separately for variation within a machine
- These scales are illustrative, not universal defaults



- Hierarchical models
 - Rats example
 - **Factory example**
 - 8 schools example
- Exchangeability
- Computational aspects

The prior on variation affects pooling

- Small τ : machine means lie close to μ , encouraging stronger pooling
- Large τ : machine means can differ more, allowing less pooling
- With few groups, the data may give little information about τ , so its prior can matter substantially
- Check prior predictive outcomes and compare plausible prior scales
- A uniform prior on τ , on τ^2 , or on $\log \tau$ gives different assumptions; “flat” does not mean neutral

BDA3, Section 5.7 discusses priors for hierarchical variation. The half-normal specification here is one modelling choice.



- Hierarchical models
 - Rats example
 - **Factory example**
 - 8 schools example
- Exchangeability
- Computational aspects

Prediction for an existing machine

In the common-variance model, distinguish the mean θ_j from a future observation $\tilde{y}_j$:

$$p(\tilde{y}_j | y) = \int \int p(\tilde{y}_j | \theta_j, \sigma) p(\theta_j, \sigma | y) d\theta_j d\sigma.$$

For each posterior draw s :

1. Take $(\theta_j^{(s)}, \sigma^{(s)})$ from the **same joint posterior draw**
2. Draw $\tilde{y}_j^{(s)} \sim \mathcal{N}(\theta_j^{(s)}, (\sigma^{(s)})^2)$
 - The $\theta_j^{(s)}$ summarize uncertainty about the machine's mean
 - The $\tilde{y}_j^{(s)}$ also include variation in a new observation



- Hierarchical models
 - Rats example
 - Factory example
 - 8 schools example
- Exchangeability
- Computational aspects

Prediction for a new similar machine

With no observations yet from the new machine,

$$p(\theta_{\text{new}} \mid y) = \int \int p(\theta_{\text{new}} \mid \mu, \tau) p(\mu, \tau \mid y) d\mu d\tau,$$

$$p(\tilde{y}_{\text{new}} \mid y) = \int \int p(\tilde{y}_{\text{new}} \mid \theta_{\text{new}}, \sigma) p(\theta_{\text{new}}, \sigma \mid y) d\theta_{\text{new}} d\sigma.$$

For each posterior draw s :

1. Take $(\mu^{(s)}, \tau^{(s)}, \sigma^{(s)})$ jointly from $p(\mu, \tau, \sigma \mid y)$
2. Draw the new mean: $\theta_{\text{new}}^{(s)} \sim \mathcal{N}(\mu^{(s)}, (\tau^{(s)})^2)$
3. For a new observation, also draw $\tilde{y}_{\text{new}}^{(s)} \sim \mathcal{N}(\theta_{\text{new}}^{(s)}, (\sigma^{(s)})^2)$

Predicting a new mean stops at step 2; predicting an observation includes step 3.



- Hierarchical models
 - Rats example
 - **Factory example**
 - 8 schools example
- Exchangeability
- Computational aspects

Pair discussion: a new machine

Discuss in pairs (1–2 minutes).

We have data from six machines and want to predict for a new, similar machine 7.

- Its mean quality, θ_7
- The quality of its next item, $\tilde{y}_7$

How do the simulation steps differ? What additional uncertainty enters the second prediction?



Pair discussion: a new machine

Discuss in pairs (1–2 minutes).

We have data from six machines and want to predict for a new, similar machine 7.

- Its mean quality, θ_7
- The quality of its next item, $\tilde{y}_7$

How do the simulation steps differ? What additional uncertainty enters the second prediction?

Answer: Both start by drawing a new machine mean using posterior draws of the population parameters. Predicting an item also adds within-machine variation:


$$\tilde{y}_7^{(s)} \sim \mathcal{N}(\theta_7^{(s)}, (\sigma^{(s)})^2).$$



- Hierarchical models
 - Rats example
 - **Factory example**
 - 8 schools example
- Exchangeability
- Computational aspects

Hierarchical normal model: factory

- Six machines with means θ_j and **own variances** σ_j^2



$$y_{ij} \mid \theta_j, \sigma_j^2 \sim \mathcal{N}(\theta_j, \sigma_j^2)$$

$$\theta_j \mid \mu, \tau \sim \mathcal{N}(\mu, \tau^2)$$


$$\sigma_j^2 \mid \nu_0, \sigma_0^2 \sim \text{Inv-}\chi^2(\nu_0, \sigma_0^2)$$



- Hierarchical models
 - Rats example
 - **Factory example**
 - 8 schools example
- Exchangeability
- Computational aspects

Hierarchical normal model: factory

- Six machines with means θ_j and **own variances** σ_j^2



$$y_{ij} \mid \theta_j, \sigma_j^2 \sim \mathcal{N}(\theta_j, \sigma_j^2)$$

$$\theta_j \mid \mu, \tau \sim \mathcal{N}(\mu, \tau^2)$$

$$\sigma_j^2 \mid \nu_0, \sigma_0^2 \sim \text{Inv-}\chi^2(\nu_0, \sigma_0^2)$$


- **Pair discussion (1–2 minutes):** What predictive distributions can we have?



- Hierarchical models
 - Rats example
 - **Factory example**
 - 8 schools example
- Exchangeability
- Computational aspects

Hierarchical normal model: factory

- Six machines with means θ_j and **own variances** σ_j^2



$$y_{ij} \mid \theta_j, \sigma_j^2 \sim \mathcal{N}(\theta_j, \sigma_j^2)$$

$$\theta_j \mid \mu, \tau \sim \mathcal{N}(\mu, \tau^2)$$

$$\sigma_j^2 \mid \nu_0, \sigma_0^2 \sim \text{Inv-}\chi^2(\nu_0, \sigma_0^2)$$

- **Pair discussion (1–2 minutes):** What predictive distributions can we have?
- Observations from existing or new similar machines



Eight schools: historical SAT coaching study

- Hierarchical models
 - Rats example
 - Factory example
 - 8 schools example
- Exchangeability
- Computational aspects
- Historical example from BDA3, Section 5.5
 - randomized coaching experiments in eight high schools
 - question: does short-term coaching improve verbal test scores?
 - estimate a coaching effect for each school



Eight schools: historical SAT coaching study

- Hierarchical models
 - Rats example
 - Factory example
 - 8 schools example
- Exchangeability
- Computational aspects

- Historical example from BDA3, Section 5.5
 - randomized coaching experiments in eight high schools
 - question: does short-term coaching improve verbal test scores?
 - estimate a coaching effect for each school
- SAT-V in this study (the historical verbal test)
 - standardized multiple-choice test
 - score range 200–800, with mean about 500 and SD about 100
 - pre-test PSAT scores in verbal (V) and mathematics (M) help adjust for students' initial performance



- Hierarchical models
 - Rats example
 - Factory example
 - 8 schools example
- Exchangeability
- Computational aspects

- Effectiveness of the SAT coaching
 - students took pre-tests PSAT-M and PSAT-V
 - some students received coaching
 - a regression in each school adjusts for pre-test scores and estimates the coaching effect y_j and its standard error σ_j
 - use a normal approximation for y_j , treating the estimated sampling variances σ_j^2 as known
 - y_j is an adjusted effect estimate, not a sample mean; σ_j is not the SD of individual students' scores



- Hierarchical models
 - Rats example
 - Factory example
 - 8 schools example
- Exchangeability
- Computational aspects

- Effectiveness of the SAT coaching
 - students took pre-tests PSAT-M and PSAT-V
 - some students received coaching
 - a regression in each school adjusts for pre-test scores and estimates the coaching effect y_j and its standard error σ_j
 - use a normal approximation for y_j , treating the estimated sampling variances σ_j^2 as known
 - y_j is an adjusted effect estimate, not a sample mean; σ_j is not the SD of individual students' scores

• Data:	School	A	B	C	D	E	F	G	H
	y_j	28	8	-3	7	-1	1	18	12
	σ_j	15	10	16	11	9	11	10	18



- Hierarchical models
 - Rats example
 - Factory example
 - 8 schools example
- Exchangeability
- Computational aspects

Hierarchical normal model for group means

- Observations are independent given θ_j ; σ^2 is known

$$y_{ij} | \theta_j \sim \mathcal{N}(\theta_j, \sigma^2), \quad i = 1, \dots, n_j; \quad j = 1, \dots, J$$

- Group j sample mean and its sampling variance

$$\bar{y}_{\cdot j} = \frac{1}{n_j} \sum_{i=1}^{n_j} y_{ij}$$

$$\text{Var}(\bar{y}_{\cdot j} | \theta_j) = \frac{\sigma^2}{n_j} = \sigma_j^2$$

SD of observations: σ ; SE of the mean: $\sigma_j = \sigma / \sqrt{n_j}$.



- Hierarchical models
 - Rats example
 - Factory example
 - 8 schools example
- Exchangeability
- Computational aspects

Hierarchical normal model for group means

- Observations are independent given θ_j ; σ^2 is known

$$y_{ij} | \theta_j \sim \mathcal{N}(\theta_j, \sigma^2), \quad i = 1, \dots, n_j; \quad j = 1, \dots, J$$

- Group j sample mean and its sampling variance

$$\bar{y}_{.j} = \frac{1}{n_j} \sum_{i=1}^{n_j} y_{ij}$$

$$\text{Var}(\bar{y}_{.j} | \theta_j) = \frac{\sigma^2}{n_j} = \sigma_j^2$$

SD of observations: σ ; SE of the mean: $\sigma_j = \sigma / \sqrt{n_j}$.

- Thus $\bar{y}_{.j} | \theta_j \sim \mathcal{N}(\theta_j, \sigma_j^2)$.
The same form applies to adjusted school effects y_j , using their reported standard errors σ_j .




- Hierarchical models
 - Rats example
 - Factory example
 - 8 schools example
- Exchangeability
- Computational aspects

Hierarchical normal model for group means

$$\theta_j | \mu, \tau \sim \mathcal{N}(\mu, \tau)$$

$$\bar{y}_{.j} | \theta_j \sim \mathcal{N}(\theta_j, \sigma_j^2)$$



Here $\mathcal{N}(\mu, \tau)$ is written in mean–SD form; τ is the between-group SD.



- Hierarchical models
 - Rats example
 - Factory example
 - 8 schools example
- Exchangeability
- Computational aspects

Normal shrinkage: completing the square

For fixed μ and $\tau > 0$, with known sampling SD $\sigma_j > 0$,

$$y_j \mid \theta_j \sim \mathcal{N}(\theta_j, \sigma_j^2), \quad \theta_j \mid \mu, \tau \sim \mathcal{N}(\mu, \tau^2).$$

Here $\mathcal{N}$ uses variance as its second argument.

$$\begin{aligned} p(\theta_j \mid \mu, \tau, y) &\propto \exp \left[-\frac{1}{2} \left\{ \frac{(y_j - \theta_j)^2}{\sigma_j^2} + \frac{(\theta_j - \mu)^2}{\tau^2} \right\} \right] \\ &\propto \exp \left[-\frac{(\theta_j - m_j)^2}{2v_j} \right], \end{aligned}$$

where completing the square gives

$$v_j = \left(\frac{1}{\sigma_j^2} + \frac{1}{\tau^2} \right)^{-1}, \quad m_j = v_j \left(\frac{y_j}{\sigma_j^2} + \frac{\mu}{\tau^2} \right).$$

Thus $\theta_j \mid \mu, \tau, y \sim \mathcal{N}(m_j, v_j)$.



Rearranging the conditional posterior mean,

$$E(\theta_j \mid \mu, \tau, y) = w_j y_j + (1 - w_j) \mu, \quad w_j = \frac{\tau^2}{\tau^2 + \sigma_j^2}.$$

- A precise group estimate (small σ_j) receives more weight
- Small between-group variation τ gives more weight to μ
- As $\tau \rightarrow 0$, the conditional mean approaches μ ; as $\tau \rightarrow \infty$, it approaches y_j
- In full Bayes, average these conditional results over $p(\mu, \tau \mid y)$



Normal shrinkage: a numerical example

Use School A's $y_j = 28$, $\sigma_j = 15$ and, for illustration, **fix** $\mu = 8$, $\tau = 5$.

$$w_j = \frac{5^2}{5^2 + 15^2} = \frac{25}{250} = 0.10,$$

$$m_j = 0.10 \cdot 28 + 0.90 \cdot 8 = 10,$$

$$v_j = \left(\frac{1}{225} + \frac{1}{25} \right)^{-1} = 22.5.$$

- The noisy school estimate receives 10% of the weight
- Conditional posterior: mean 10, SD $\sqrt{22.5} \approx 4.74$
- This illustrates fixed hyperparameters; the full eight-schools posterior also averages over μ, τ



Pair discussion: precision and shrinkage

Discuss in pairs (1–2 minutes).

Two schools have the same estimated effect, different from the population mean μ , but different standard errors:

$$\sigma_1 = 5, \quad \sigma_2 = 20.$$

Hold μ and $\tau > 0$ fixed.

Which school's estimate is shrunk more toward μ ? Why?

- Hierarchical models
 - Rats example
 - Factory example
 - 8 schools example
- Exchangeability
- Computational aspects



Pair discussion: precision and shrinkage

Discuss in pairs (1–2 minutes).

Two schools have the same estimated effect, different from the population mean μ , but different standard errors:

$$\sigma_1 = 5, \quad \sigma_2 = 20.$$

Hold μ and $\tau > 0$ fixed.

Which school's estimate is shrunk more toward μ ? Why?

Answer: School 2. Its larger standard error means less precise evidence, so its estimate receives less weight and μ receives more weight.

- Hierarchical models
 - Rats example
 - Factory example
 - 8 schools example
- Exchangeability
- Computational aspects




- Hierarchical models
 - Rats example
 - Factory example
 - 8 schools example
- Exchangeability
- Computational aspects

Eight schools: pooling assumptions

How does the posterior for School A (red) change?

Separate model



No pooling: each school's effect uses its own data. School A's estimate is high but uncertain.

The hierarchical plots follow BDA3, Section 5.5, with flat independent priors on μ and $\tau > 0$, rather than the illustrative half-normal prior above.




- Hierarchical models
 - Rats example
 - Factory example
 - 8 schools example
- Exchangeability
- Computational aspects

Eight schools: pooling assumptions

How does the posterior for School A (red) change?

Pooled model



Complete pooling: all schools share a common effect. School A's posterior is identical to the others.

The hierarchical plots follow BDA3, Section 5.5, with flat independent priors on μ and $\tau > 0$, rather than the illustrative half-normal prior above.




- Hierarchical models
 - Rats example
 - Factory example
 - 8 schools example
- Exchangeability
- Computational aspects

Eight schools: pooling assumptions

How does the posterior for School A (red) change?

Hierarchical model



Partial pooling: the marginal posteriors $p(\theta_j | y)$ share information while allowing different effects.


The hierarchical plots follow BDA3, Section 5.5, with flat independent priors on μ and $\tau > 0$, rather than the illustrative half-normal prior above.



- Hierarchical models
 - Rats example
 - Factory example
 - 8 schools example
- Exchangeability
- Computational aspects

Eight schools: uncertainty about variation

Does the peak near zero establish that all effects are equal?




- This is $p(\tau | y)$, with μ and the school effects integrated out
- A range of between-school SDs remains plausible
- Full Bayes averages over this distribution; fixing τ at its mode would impose much stronger pooling



Eight schools: conditional posterior means

What happens to the school estimates as τ changes?

Conditional means $E[\theta_j|\tau, y]$




- Curves show $E(\theta_j | \tau, y)$: τ is fixed and μ is integrated out
- Small τ : means come together near the pooled estimate, 7.7
- Large τ : means approach the separate estimates y_j



Eight schools: uncertainty given τ

Does small τ make the common effect precisely known?

Conditional standard deviations $\text{sd}[\theta_j|\tau, y]$



- Curves show $\text{sd}(\theta_j \mid \tau, y)$, with uncertainty in μ included
- As $\tau \rightarrow 0$, effects coincide, but their common value is uncertain:

$$\text{sd}(\theta_j \mid \tau, y) \longrightarrow \left(\sum_{k=1}^J \frac{1}{\sigma_k^2} \right)^{-1/2} \approx 4.1.$$



UPPSALA
UNIVERSITET

- Hierarchical models
 - Rats example
 - Factory example
 - 8 schools example
- Exchangeability
- Computational aspects

Section 2

Exchangeability



- Hierarchical models
 - Rats example
 - Factory example
 - 8 schools example
- Exchangeability
- Computational aspects

- Motivates a **symmetric joint model**
 - for observations whose labels carry no relevant information
 - for group parameters that we treat symmetrically a priori
- Less restrictive than **IID** (independent and identically distributed)
- IID implies exchangeability, but exchangeability need not imply independence
- Independence alone is not enough: the distributions may differ



- Hierarchical models
 - Rats example
 - Factory example
 - 8 schools example
- Exchangeability
- Computational aspects

Exchangeability

- *Exchangeability*

Random variables $\theta_1, \dots, \theta_J$ (or $y_1, \dots, y_J$) are **exchangeable** if the joint distribution p is invariant to the permutation of indices

i.e

$$p(\theta_1, \theta_2, \theta_3) = p(\theta_{\pi(1)}, \theta_{\pi(2)}, \theta_{\pi(3)})$$

for any permutation π of the indices. E.g.

$$p(\theta_1, \theta_2, \theta_3) = p(\theta_2, \theta_3, \theta_1)$$



- Hierarchical models
 - Rats example
 - Factory example
 - 8 schools example
- Exchangeability
- Computational aspects

Exchangeability

- *Exchangeability*

Random variables $\theta_1, \dots, \theta_J$ (or $y_1, \dots, y_J$) are **exchangeable** if the joint distribution p is invariant to the permutation of indices

i.e

$$p(\theta_1, \theta_2, \theta_3) = p(\theta_{\pi(1)}, \theta_{\pi(2)}, \theta_{\pi(3)})$$

for any permutation π of the indices. E.g.

$$p(\theta_1, \theta_2, \theta_3) = p(\theta_2, \theta_3, \theta_1)$$

Can we come up with a situation where this doesn't hold?



- Hierarchical models
 - Rats example
 - Factory example
 - 8 schools example
- Exchangeability
- Computational aspects

Exchangeability

- *Exchangeability*

Random variables $\theta_1, \dots, \theta_J$ (or $y_1, \dots, y_J$) are **exchangeable** if the joint distribution p is invariant to the permutation of indices

i.e

$$p(\theta_1, \theta_2, \theta_3) = p(\theta_{\pi(1)}, \theta_{\pi(2)}, \theta_{\pi(3)})$$

for any permutation π of the indices. E.g.

$$p(\theta_1, \theta_2, \theta_3) = p(\theta_2, \theta_3, \theta_1)$$

Can we come up with a situation where this doesn't hold?
E.g. when we have a trend over indices (like in time series data)



- Hierarchical models
 - Rats example
 - Factory example
 - 8 schools example
- Exchangeability
- Computational aspects

Exchangeability

- *Exchangeability*

Random variables $\theta_1, \dots, \theta_J$ (or $y_1, \dots, y_J$) are **exchangeable** if the joint distribution p is invariant to the permutation of indices

i.e

$$p(\theta_1, \theta_2, \theta_3) = p(\theta_{\pi(1)}, \theta_{\pi(2)}, \theta_{\pi(3)})$$

for any permutation π of the indices. E.g.

$$p(\theta_1, \theta_2, \theta_3) = p(\theta_2, \theta_3, \theta_1)$$

Can we come up with a situation where this doesn't hold?
E.g. when we have a trend over indices (like in time series data)

- Exchangeability implies symmetry:
If there is no information which can be used *a priori* to separate θ_j from each other, we can assume exchangeability. ("Ignorance implies exchangeability")



- Hierarchical models
 - Rats example
 - Factory example
 - 8 schools example
- Exchangeability
- Computational aspects

Exchangeability: de Finetti's theorem

- **Bernoulli case:** let $(X_i)_{i=1}^{\infty}$ be an infinite exchangeable sequence with $X_i \in \{0, 1\}$
- There is a random $\theta \in [0, 1]$ such that the X_i are **conditionally IID**:

$$X_i \mid \theta \stackrel{\text{iid}}{\sim} \text{Bernoulli}(\theta)$$

- If θ has a density $p(\theta)$, then for every finite J ,

$$p(x_1, \dots, x_J) = \int_0^1 \left[\prod_{i=1}^J \theta^{x_i} (1 - \theta)^{1-x_i} \right] p(\theta) d\theta$$

- The theorem does not choose the mixing distribution: a Beta prior is a modelling choice, not a consequence



Coin flips: exchangeable but dependent

Let $\theta \sim \text{Beta}(1, 1)$ and, given θ , let the flips be IID Bernoulli(θ). Write H_i for heads on flip i .

$$\Pr(H_1) = \mathbb{E}(\theta) = \frac{1}{2}.$$

- Hierarchical models
 - Rats example
 - Factory example
 - 8 schools example
- Exchangeability
- Computational aspects



Coin flips: exchangeable but dependent

Let $\theta \sim \text{Beta}(1, 1)$ and, given θ , let the flips be IID Bernoulli(θ). Write H_i for heads on flip i .

$$\Pr(H_1) = \mathbb{E}(\theta) = \frac{1}{2}.$$

After observing H_1 ,

$$\theta \mid H_1 \sim \text{Beta}(2, 1), \quad \Pr(H_2 \mid H_1) = \mathbb{E}(\theta \mid H_1) = \frac{2}{3}.$$

- Marginally $\Pr(H_2) = 1/2$, so the flips are dependent
- Their joint distribution remains unchanged when we swap the flips
- Observing heads updates the shared unknown bias; one toss does not physically affect the next



- Hierarchical models
 - Rats example
 - Factory example
 - 8 schools example
- Exchangeability
- Computational aspects

Pair discussion: learning about the coin

Discuss in pairs (1–2 minutes).

The coin's bias has prior $\theta \sim \text{Beta}(1, 1)$, and the flips are independent given θ .

Why does observing heads change our prediction for the next flip? Does this contradict conditional independence?



Pair discussion: learning about the coin

Discuss in pairs (1–2 minutes).

The coin's bias has prior $\theta \sim \text{Beta}(1, 1)$, and the flips are independent given θ .

Why does observing heads change our prediction for the next flip? Does this contradict conditional independence?

Answer: We learn about the shared unknown bias. After heads, $\theta \mid H_1 \sim \text{Beta}(2, 1)$, changing the predictive probability from $1/2$ to $2/3$.

The coin has not changed. Independence holds when θ is fixed; averaging over an uncertain θ creates dependence.

- Hierarchical models
 - Rats example
 - Factory example
 - 8 schools example
- Exchangeability
- Computational aspects



- Hierarchical models
 - Rats example
 - Factory example
 - 8 schools example
- Exchangeability
- Computational aspects

- Exchangeability does not mean that the results of the experiments could not be different
 - e.g. if we know that the experiments have been in two different laboratories, and we know that the other laboratory has better conditions for the rats, but we do not know which experiments have been made in which laboratory
 - a priori experiments are exchangeable
 - model could have unknown parameter for the laboratory with a conditional prior for rats assumed to come from the same place (clustering model)



UPPSALA
UNIVERSITET

Exchangeability and additional information

- Bioassay: doses x_j and group sizes n_j are fixed design information; y_j counts deaths

- Hierarchical models
 - Rats example
 - Factory example
 - 8 schools example
- Exchangeability
- Computational aspects



- Hierarchical models
 - Rats example
 - Factory example
 - 8 schools example
- Exchangeability
- Computational aspects

Exchangeability and additional information

- Bioassay: doses x_j and group sizes n_j are fixed design information; y_j counts deaths
- Given α, β , the counts are independent, but generally not identically distributed:

$$y_j \mid \alpha, \beta, n_j, x_j \sim \text{Binomial}(n_j, \pi_j), \quad j = 1, \dots, J$$

where $\text{logit}(\pi_j) = \alpha + \beta x_j$.

$$p(\alpha, \beta \mid y, n, x) \propto p(\alpha, \beta) \prod_{j=1}^J p(y_j \mid \alpha, \beta, n_j, x_j)$$



- Hierarchical models
 - Rats example
 - Factory example
 - 8 schools example
- Exchangeability
- Computational aspects

Exchangeability and additional information

- Bioassay: doses x_j and group sizes n_j are fixed design information; y_j counts deaths
- Given α, β , the counts are independent, but generally not identically distributed:

$$y_j \mid \alpha, \beta, n_j, x_j \sim \text{Binomial}(n_j, \pi_j), \quad j = 1, \dots, J$$

where $\text{logit}(\pi_j) = \alpha + \beta x_j$.

$$p(\alpha, \beta \mid y, n, x) \propto p(\alpha, \beta) \prod_{j=1}^J p(y_j \mid \alpha, \beta, n_j, x_j)$$

- Reordering complete records (x_j, n_j, y_j) leaves the likelihood unchanged; swapping only counts between doses generally changes it



UPPSALA
UNIVERSITET

Hierarchical exchangeability

- Hierarchical models
 - Rats example
 - Factory example
 - 8 schools example
 - Exchangeability
 - Computational aspects
- Rat example: experiments $j = 1, \dots, J$, each with n_j rats



UPPSALA
UNIVERSITET

Hierarchical exchangeability

- Hierarchical models
 - Rats example
 - Factory example
 - 8 schools example
 - Exchangeability
 - Computational aspects
- Rat example: experiments $j = 1, \dots, J$, each with n_j rats
 - Within experiment j , outcomes are IID Bernoulli(θ_j) given θ_j ; they need not be exchangeable across experiments



UPPSALA
UNIVERSITET

Hierarchical exchangeability

- Hierarchical models
 - Rats example
 - Factory example
 - 8 schools example
- Exchangeability
- Computational aspects

- Rat example: experiments $j = 1, \dots, J$, each with n_j rats
- Within experiment j , outcomes are IID Bernoulli(θ_j) given θ_j ; they need not be exchangeable across experiments
- The group probabilities θ_j are modelled exchangeably a priori



UPPSALA
UNIVERSITET

Hierarchical exchangeability

- Hierarchical models
 - Rats example
 - Factory example
 - 8 schools example
- Exchangeability
- Computational aspects

- Rat example: experiments $j = 1, \dots, J$, each with n_j rats
- Within experiment j , outcomes are IID Bernoulli(θ_j) given θ_j ; they need not be exchangeable across experiments
- The group probabilities θ_j are modelled exchangeably a priori
- Counts $y_j \mid \theta_j, n_j \sim \text{Binomial}(n_j, \theta_j)$ need not be exchangeable when n_j differs
- This assumes comparable experiments; grouping alone does not justify a shared population distribution



- Hierarchical models
 - Rats example
 - Factory example
 - 8 schools example
- Exchangeability
- Computational aspects

- **Conditional exchangeability**

- Responses with the same covariates may be modelled exchangeably; different covariates may imply different distributions
- For fixed covariates, use a conditional model $p(y \mid x)$; adding x_i does not make (x_i, y_i) exchangeable



- Hierarchical models
 - Rats example
 - Factory example
 - 8 schools example
- Exchangeability
- Computational aspects

- **Conditional exchangeability**

- Responses with the same covariates may be modelled exchangeably; different covariates may imply different distributions
- For fixed covariates, use a conditional model $p(y \mid x)$; adding x_i does not make (x_i, y_i) exchangeable

- **Partial exchangeability**

- Symmetry within specified groups, not under arbitrary exchanges across groups
- The models here additionally assume exchangeable group parameters, possibly given group-level covariates



UPPSALA
UNIVERSITET

- Hierarchical models
 - Rats example
 - Factory example
 - 8 schools example
- Exchangeability
- Computational aspects

Section 3

Computational aspects



- Consider the **prior hierarchy**, before observing data:

$$\tau \sim p(\tau), \quad \tau > 0, \quad \theta_j \mid \tau \sim \mathcal{N}(0, \tau) \quad (\text{SD form})$$

- Small τ : group effects lie near 0; large τ : group effects are more spread out
- If plausible τ values range from near zero to large values, the joint prior has a narrow neck and a broad region: a funnel

- Hierarchical models
 - Rats example
 - Factory example
 - 8 schools example
- Exchangeability
- Computational aspects



- Hierarchical models
 - Rats example
 - Factory example
 - 8 schools example
- Exchangeability
- Computational aspects

Funnel geometry: prior and posterior

- Consider the **prior hierarchy**, before observing data:

$$\tau \sim p(\tau), \quad \tau > 0, \quad \theta_j \mid \tau \sim \mathcal{N}(0, \tau) \quad (\text{SD form})$$


- Small τ : group effects lie near 0; large τ : group effects are more spread out
- If plausible τ values range from near zero to large values, the joint prior has a narrow neck and a broad region: a funnel
- With weak information about group effects, the **posterior can retain** this funnel geometry
- Not every hierarchical posterior is funnel-shaped: the likelihood also matters



- Hierarchical models
 - Rats example
 - Factory example
 - 8 schools example
- Exchangeability
- Computational aspects

Why is a funnel problematic?

- Changing curvature makes it hard for HMC to explore both the narrow and wide regions efficiently
- Post-warmup divergences warn of potentially unreliable exploration and biased posterior estimates
- Good $\hat{R}$ and ESS do not cancel divergence warnings



Betancourt (2020) [Stan diagnostics](#) [HMC demo](#)

In this figure, ϕ denotes one group effect θ_j .




- Hierarchical models
 - Rats example
 - Factory example
 - 8 schools example
- Exchangeability
- Computational aspects

Handling the funnel

- Raising `adapt_delta` targets smaller steps: it can help, but does not guarantee reliable inference
- **Centered**: often better with strong group information

$$\theta_j \sim N(\mu, \tau) \quad (\text{SD form})$$





- Hierarchical models
 - Rats example
 - Factory example
 - 8 schools example
- Exchangeability
- Computational aspects


Handling the funnel

- Raising `adapt_delta` targets smaller steps: it can help, but does not guarantee reliable inference
- **Centered:** often better with strong group information

$$\theta_j \sim N(\mu, \tau) \quad (\text{SD form})$$

- **Non-centered:** often helps weakly informed group effects

$$\eta_j \sim N(0, 1), \quad \theta_j = \mu + \tau \eta_j$$





- Hierarchical models
 - Rats example
 - Factory example
 - 8 schools example
- Exchangeability
- Computational aspects

Why non-centering preserves the model

Start with $\eta_j \stackrel{\text{iid}}{\sim} \mathcal{N}(0, 1)$, independent of the hyperparameters a priori, and set

$$\theta_j = \mu + \tau \eta_j, \quad \tau > 0.$$

Conditional on μ, τ , an affine transformation of a normal variable is normal, with

$$\begin{aligned} \mathbb{E}(\theta_j \mid \mu, \tau) &= \mu + \tau \mathbb{E}(\eta_j) = \mu, \\ \text{Var}(\theta_j \mid \mu, \tau) &= \tau^2 \text{Var}(\eta_j) = \tau^2. \end{aligned}$$

Thus the θ_j are conditionally independent normals with mean μ and SD τ , as in the centered hierarchy.

Keep the same hyperpriors and likelihood, evaluated at $\theta_j = \mu + \tau \eta_j$. Only the sampling coordinates change.